Frex: Dependently Typed Algebraic Simplification

GUILLAUME ALLAIS, University of Strathclyde, United Kingdom
EDWIN BRADY, University of St. Andrews, United Kingdom
NATHAN CORBYN, University of Oxford, United Kingdom
OHAD KAMMAR, University of Edinburgh, United Kingdom
JEREMY YALLOP, University of Cambridge, United Kingdom

We present a new design for an algebraic simplification library structured around concepts from universal algebra: theories, models, homomorphisms, and universal properties of free algebras and free extensions of algebras. The library's dependently typed interface guarantees that both built-in and user-defined simplification modules are terminating, sound, and complete with respect to a well-specified class of equations. We have implemented the design in the Idris 2 and Agda dependently typed programming languages and shown that it supports modular extension to new theories, proof extraction and certification, goal extraction via reflection, and interactive development.

CCS Concepts: • **Theory of computation** → **Type theory**; **Constructive mathematics**; **Equational logic and rewriting**; **Automated reasoning**; *Categorical semantics*; *Algebraic semantics*; • **Software and its engineering** → Formal software verification; Functional languages; • **Mathematics of computing** → Solvers; • **Computing methodologies** → Representation of polynomials.

Additional Key Words and Phrases: dependent types, frex, free extension, mathematically structured programming, universal algebra, algebraic simplification, homomorphism, universal property

1 Introduction

Algebraic simplification involves using algebraic laws to normalise expressions with unknowns. For example, the commutative monoid axioms—associativity, neutrality, and commutativity—over integers with addition serve to simplify the left hand expression to the right hand expression below:

$$-6 + (x + 3) + (y + x) \xmapsto{\text{simplify}} -3 + 2x + y$$

Many application domains make use of this kind of simplification. For example, algebraic simplification is often a useful first step in a program optimiser, to avoid the need to analyse and transform distinct but equivalent programs. The present work focuses on another application domain: theorem provers and programming languages based on type theory. In these systems, users often need to prove to a type checker that terms are equivalent, but constructing the proofs often involves rote algebraic simplification steps that users resent having to produce manually.

Authors' Contact Information: Guillaume Allais, University of Strathclyde, Glasgow, United Kingdom, guillaume.allais@ens-lyon.org; Edwin Brady, University of St. Andrews, St. Andrews, United Kingdom, ecb10@st-andrews.ac.uk; Nathan Corbyn, University of Oxford, Oxford, United Kingdom, nathan.corbyn@cs.ox.ac.uk; Ohad Kammar, University of Edinburgh, Edinburgh, United Kingdom, ohad.kammar@ed.ac.uk; Jeremy Yallop, University of Cambridge, Cambridge, United Kingdom, jeremy.yallop@cl.cam.ac.uk.

To free users from the need to construct rote algebraic proofs, dependently typed languages and their ecosystems often include simplifiers for common algebraic structures such as monoids, semi-rings, and rings. With these simplifiers, users need only establish the structures' axioms, such as neutrality, associativity and commutativity, and can then call the simplifiers to discharge rote simplification steps. Implementation strategies for simplifiers vary: some, such as the Rocq ring solver [Barras et al. 2021], use tactics to simplify algebraic terms in typing goals, while others, such as the Agda ring solver [Kidney 2019], use proof-by-reflection to construct propositions to discharge equations. Some simplifiers, such as the Rocq ring solver, group together several algebraic structures, while others, such as the system by Grégoire and Mahboubi [2005], generalise several distinct structures to one structure. The state-of-the-art are standalone simplifiers that, through heuristics and long-term development, can deal with common cases.

1.1 Representation Theorems

Universal algebra has a long tradition concerning algebraic simplification under the collective name 'representation theorems'. Each such representation theorem characterises canonical representatives of algebraic expressions in terms of (typically inductive) constructions such as reduced-words and formal polynomials. Such characterisations often reuse existing representation theorems of simpler algebraic structures or familiar algebraic structures such as the integers or the natural numbers.

The present work makes use of two kinds of representation theorems. For a free algebra (abbreviated *fral*), a representation theorem amounts to an algebraic structure that interprets the algebraic operations to produce a single canonical representative for all input terms that are equivalent according to the algebraic laws. For a free extension (abbreviated *frex*), a representation theorem chooses a canonical representative using the algebraic laws while also evaluating concrete elements. Free algebras and free extensions are related: each free extension is also a free algebra of a theory specialised to a concrete algebra by adding axioms over the concrete elements.

For example, for commutative monoids the fral representation theorem states that the free commutative monoid over n variables is represented by the set of n-tuples of naturals $\mathbb{N}^n$. We simplify with this representation by evaluating a term in the fral and then reifying it back as a term:

$$-6+(x+3)+(y+x) \xmapsto{\text{evaluate}(x_0+(x_2+x_1)+(x_3+x_2))} (1,1,2,1) \xmapsto{\text{reify}(x_0\mapsto-6,x_1\mapsto3,x_2\mapsto x,x_3\mapsto y)} -6+3+2x+y$$

In the fral there is no concept of concrete elements, so fral simplification must treat constants such as -6 and 3 in the same way as variables such as x and y, as abstract and distinct indeterminates.

The frex representation theorem for commutative monoids states that the free extension of a commutative monoid over C by n variables is represented by the set $C \times \mathbb{N}^n$, where the element $(c, a_1, \ldots, a_n)$ represents the expression $c + a_1 x_1 + \ldots + a_n x_n$. Simplifying a term using the frex representation involves evaluating a term in the frex and then reifying it back as a term:

$$-6 + (x + 3) + (y + x) \xmapsto{\text{evaluate}_\mathbb{Z}} (-3, 2, 1) \xmapsto{\text{reify}} -3 + 2x + y$$

The frex representation distinguishes variables from concrete elements, gathering the latter together and evaluating them using the operations of the concrete commutative monoid in use.

Both kinds of representation theorem allow us to choose a unique syntactic representative, i.e., a normal form. Such normalization by evaluation and then reification also applies to more sophisticated notions of algebra that include the equational theories of λ-calculi, and is familiar in those settings as *normalization-by-evaluation* [Berger and Schwichtenberg 1991]. Its first systematic applications were in the formal study of various simply typed calculi [Altenkirch et al. 2001, 1995; Čubrić et al. 1998] and category theoretic constructions [Beylin and Dybjer 1996]. It has served as a conversion-checking technique during the type-checking of dependently typed calculi from their

inception [Martin-Löf 1975], gaining adoption after the seminal works of Abel et al. [2007a,b], even for sophisticated calculi [Abel et al. 2017; Hu et al. 2023; Sterling and Angiuli 2021].

This article describes an extensible dependently typed library for algebraic simplifiers based on fral and frex representation theorems, building on our previous work on free extensions for partial evaluation [Yallop et al. 2018]. Current implementations of algebraic simplifiers, even in dependently typed settings, are restricted to implementing the computational representation— i.e. the data-structures needed for the normal form together with the normalisation-by-evaluation algorithm—alongside a formalisation of the soundness proofs. This work investigates what can be gained by encoding, in addition, the full meta-theory of these representation theorems, including generic representations of theories, their algebras and algebra homomorphisms, and the universal properties of the fral and frex. For simplicity of development and exposition we apply this generic machinery to a handful of familiar monoid varieties. Moreover, we ensure all of these concepts remain computational by avoiding the temptation of postulating axioms that could hinder reduction of closed terms. This last task is challenging to combine with interactive performance, as formalising universal properties tends to produce large terms that slow type-checkers [Gross et al. 2014].

1.2 Paper Outline and Contributions

Section 2 present background material: a review of the mathematical foundation for Frex, and a brief Idris2 tutorial that reviews setoid-based equational reasoning.

Sections 3 and 4 present our central contribution, a fundamentally new approach to building algebraic simplifiers. The standard existing approach is to write an ad-hoc simplifier for some particular algebraic structure such as rings. We take a different approach, teaching the implementation the basic concepts of universal algebra — signatures, theories and models, homomorphisms, and universal properties (Section 3) — then build a completely generic solver based on free algebras and free extensions that can be instantiated with a particular algebra to discharge concrete proof obligations (Section 4). This new approach is inherently modular and extensible, and delivers solvers that are sound and complete by construction. We have created implementations of our design in two dependently typed languages: Fragment in Agda, and Frex in Idris2[1].

Section 5 explains the completeness guarantees of the library, and covers proof extraction, simplification, pretty-printing and certification. The technical Sections 4 and 5 are aimed at library designers, and may be skimmed or skipped at first reading.

Section 6 considers a natural question: can one use reflection to invoke Frex automatically? The answer is a qualified 'yes', requiring much library-developer effort, but leading to real advantages in Agda and limited advantages in Idris2.

Section 7 reports some supplementary evaluation of Frex. The key properties of our design are guaranteed by the type theories of the languages in which we realise it: it delivers sound and complete solvers in a completely generic way, with support for proof extraction, certification, etc. However, the practical questions of usability and viability for interactive development cannot be established by theorems and so we have also carried out some experiments. These experiments focus on the varieties of monoids that also serve as our running example, and establish that the generic solver is comfortably fast enough for interactive use, and can be extended with new algebras in a modular way and without enormous effort.

Section 8 discusses system design issues that we encountered with Frex, and Sections 9 and 10 conclude with related and further work.

Appendices A–C have more information about Frex's 9,500-line Idris2 codebase, and example code extraction.

---

[1]Available here: https://github.com/frex-project/agda-fragment and here: https://github.com/frex-project/idris-frex

2 Mathematical Overview

Universal algebra concerns the generic description of algebraic structures and relationships between them [Burris and Sankappanavar 1981, Chp. II]. We summarise the concepts relevant to Frex.

2.1 Presentations of Algebraic Structures

A *finitary signature* $\Sigma = (\mathsf{Op}\,\Sigma, \mathtt{arity})$ consists of a set $\mathsf{Op}\,\Sigma$ of *operators*, and an assignment $\mathtt{arity}$: $\mathsf{Op}\,\Sigma \to \mathtt{Nat}$ of a natural number to each operator called its *arity*. For example, the signature $\Sigma_{(2,0)}$ often used for monoids has two operators $\mathsf{Op}(\Sigma_{(2,0)}) := \{(+), 0\}$ with respective arities 2 and 0. It is common to use a succinct notation that groups operators and their arities: $\Sigma_{(2,0)} := \{(+) : 2, 0 : 0\}$.

Signatures determine an algebraic language, whose semantic models are *algebras*. An *algebra* $\mathsf{A} = (\mathsf{U}\,\mathsf{A}, \mathsf{A}[\![-]\!])$ for a signature Σ consists of a set $\mathsf{U}\,\mathsf{A}$ called the *carrier* and an assignment $\mathsf{A}[\![f]\!]$: $(\mathsf{U}\,\mathsf{A})^n \to \mathsf{U}\,\mathsf{A}$ of an *n-ary operation* over this carrier for every *n*-ary operator $f : n$ in Σ. Continuing the example above, $\Sigma_{(2,0)}$-algebras amount to triples $(X, [\![(+)]\!] : X^2 \to X, [\![0]\!] \in X)$. There may be many different algebras for a given signature and carrier set: for instance, we can equip the natural numbers $\mathbb{N}$ with the $\Sigma_{(2,0)}$-algebra structures of arithmetic addition $(\mathbb{N}, (+), 0)$, arithmetic multiplication $(\mathbb{N}, (\cdot), 1)$, or $(\mathbb{N}, \max, 0)$. Examples abound: subtraction over the integers has a $\Sigma_{(2,0)}$-algebra $(\mathbb{Z}, (-), 0)$; $n \times n$ matrices over $\mathbb{N}$ have $\Sigma_{(2,0)}$-algebra structures given by matrix addition and multiplication with the zero and identity matrix respectively, and so on.

Each signature determines a language consisting of *terms*. Given a set X of variables, the Σ-*terms over* X are given inductively as either a variable in X or an application $f(t_1, \ldots, t_n)$ of an operator $f : n$ from Σ to n terms over X. The primary role of terms is to designate *equations in context* $\mathsf{X} \vdash t = s$, i.e. triples consisting of a set X of variables and two terms in context X. For example, the associativity equation, expressed over the signature $\Sigma_{(2,0)}$, is $x, y, z \vdash x + (y + z) = (x + y) + z$.

An *environment* for a context (=set) X in an algebra A is a function $\mathsf{e} : \mathsf{X} \to \mathsf{U}\,\mathsf{A}$. An algebra A determines, for each term in context $\mathsf{X} \vdash t$, an *interpretation* function $\mathsf{A}[\![t]\!] : (\mathsf{U}\,\mathsf{A})^{\mathsf{X}} \to \mathsf{U}\,\mathsf{A}$. Given an environment e, $\mathsf{A}[\![t]\!]$ interprets each variable in the environment $\mathsf{A}[\![\mathsf{x}]\!]\,\mathsf{e} := \mathsf{e}\,\mathsf{x}$, and structurally interprets each operator application: $\mathsf{A}[\![f(t_1, \ldots, t_n)]\!]\,\mathsf{e} := \mathsf{A}[\![f]\!]\,(\mathsf{A}[\![t_1]\!]\,\mathsf{e}, \ldots, \mathsf{A}[\![t_n]\!]\,\mathsf{e})$. For example, the interpretation of the left-hand-side (LHS) of the associativity axiom in the $\Sigma_{(2,0)}$-algebra $(\mathbb{N}, \max, 0)$ given the environment $\{x \mapsto 5, y \mapsto 3, z \mapsto 8\}$ is $\max(5, \max(3, 8)) = 8$.

We say that an equation is *valid* in an algebra A, writing $\mathsf{A} \models (\mathsf{X} \vdash t = s)$, when $\mathsf{A}[\![t]\!]\,\mathsf{e} = \mathsf{A}[\![s]\!]\,\mathsf{e}$ for all environments $\mathsf{e} : \mathsf{X} \to \mathsf{U}\,\mathsf{A}$. For example, the $\Sigma_{(2,0)}$-algebras presented so far validate the associativity axiom, but interpreting the binary operation as subtraction over the integers $(\mathbb{Z}, (-), 0)$ does not: taking the environment $\{x \mapsto 0, y \mapsto 0, z \mapsto 1\}$, we have $0 - (0 - 1) = 1 \neq -1 = (0 - 0) - 1$.

A *presentation* of an algebraic theory $\mathcal{T} = (\Sigma_{\mathcal{T}}, \mathcal{T}.\mathtt{Axiom})$ consists of a signature $\Sigma_{\mathcal{T}}$ and a set $\mathcal{T}.\mathtt{Axiom}$ of $\Sigma_{\mathcal{T}}$-equations in context, which we call *axioms*. A $\mathcal{T}$-*model* A is a $\Sigma_{\mathcal{T}}$-algebra A validating all $\mathcal{T}$-axioms. For example, the axioms of the **Monoid** presentation consist of associativity and neutrality $(x \vdash x + 0 = x, 0 + x = x)$ over the signature $\Sigma_{(2,0)}$. The axioms for the **CommutativeMonoid** presentation additionally include commutativity $x, y \vdash y + x = x + y$. We can now generically manipulate classes of algebraic structures using these concepts, while generalising the usual examples: **Monoid**-models are monoids, **CommutativeMonoid**-models are commutative monoids, etc.

2.2 Homomorphisms, Free Models/Algebras, and Free Extensions

For a presentation $\mathcal{T}$, we consider two classes of simplification problem of interest, which we call *purely abstract* and *partially concrete*, respectively. In the purely abstract case, the input to the problem consists of a set of variables X and a $\Sigma_{\mathcal{T}}$-term over X. Usually, $\mathsf{X} = \mathtt{Fin}\,n$ is a set of De Bruijn levels. In the partially concrete case, the input instead consists of a $\mathcal{T}$-model A, a set of variables X and a $\Sigma_{\mathcal{T}}$-term over the disjoint union $\mathsf{U}\,\mathsf{A} \uplus \mathsf{X}$. Here, we do not treat the 'variables' in the term

coming from the inclusion of $U \, A$ as variables in the usual sense. Indeed, in both cases, the goal of the simplification problem is to find a representative modulo the presentation's axioms and the rules of deduction. However, in the partially concrete case, simplification must also fold the elements of $U \, A$ using the operations of A. When the set X is empty, simplification and evaluation should coincide.

The foregoing description is imprecise, but can be made precise using the notions of free algebra and free extension, whose definitions are formulated in terms of the unique existence of certain structure-preserving maps. We dedicate the remainder of this section to introducing these concepts, drawing connection to the frex and fral simplification examples from page 2. We start with the fral:

$$-6+(x+3)+(y+x) \xmapsto{\text{evaluate}(x_0+(x_2+x_1)+(x_3+x_2))} (1,1,2,1) \xmapsto{\text{reify}\begin{pmatrix}x_0 \mapsto -6, x_1 \mapsto 3, \\ x_2 \mapsto x, x_3 \mapsto y\end{pmatrix}} -6+3+2x+y \qquad (1)$$

Let $\mathcal{T}$ be a presentation and X a set of variables. We define a $\mathcal{T}$-*model* $a = (a.\text{Model}, \text{Env } a)$ *over* X to be a $\mathcal{T}$-model $a.\text{Model}$ equipped with an X-environment in this algebra, i.e. a function $e : X \to U(a.\text{Model})$. This concept formalises the inputs to the fral simplification process from (1). The starting term is the expression $-6 + (x + 3) + (y + x)$ in the meta-language involving the meta-level variables x and y. The argument $x_0 + (x_2 + x_1) + (x_3 + x_2)$ in the label on the left arrow is a $\Sigma_{(2,0)}$-term with object-level variables from $X := \{x_0, \dots, x_3\}$. The argument to the label on the right arrow, $(x_0 \mapsto -6, x_1 \mapsto 3, x_2 \mapsto x, x_3 \mapsto y)$, is an X-environment mapping the object-level variables $x_i \in X$ to expressions in the meta-level such as the constants -6 and 3 and the meta-level variable y. To the fral simplification process these three expressions have the same status. The frex simplification process we describe later uses the concrete nature of -6 and 3, simplifying the meta-level expression further through evaluation.

Let A, B be Σ-algebras for a signature Σ. A *homomorphism* $h : A \to B$ of Σ-algebras is a semantics-preserving function $h : U \, A \to U \, B$ between their carriers. Explicitly, for all operators $f : n$ in Σ and elements $a_1, \dots, a_n$ in $U \, A$, we have: $h(A [\![f]\!](a_1, \dots, a_n)) = B [\![f]\!](h \, a_1, \dots, h \, a_n)$. A homomorphism between $\mathcal{T}$-models is a homomorphism between the underlying signature algebras. For example, the list-length function is a homomorphism from the monoid of concatenation over lists to the monoid of addition over naturals: $\text{length} : (\text{List} X, (++), []) \to (\mathbb{N}, (+), 0)$.

A *morphism* $h : a \to b$ of $\mathcal{T}$-models over X is a $\mathcal{T}$-model homomorphism that moreover makes the diagram on the right commute. A *free $\mathcal{T}$-model over* X is then a $\mathcal{T}$-model over X from which there is a unique such morphism to every other $\mathcal{T}$-model over X. This existence-and-uniqueness property is called the *universal property* of free $\mathcal{T}$-models. For example, the free commutative monoid over $X = \{x_0, \dots, x_3\}$ is the $\Sigma_{(2,0)}$-algebra over $\mathbb{N}^4$ given by componentwise arithmetic addition, and equipped with the X-environment sending x_0 to $(1, 0, 0, 0)$, etc. The unique morphism out of this algebra into any commutative monoid A, equipped with an environment e, sends $(a_0, \dots, a_3)$ to $a_0(e \, x_0) + \dots + a_3(e \, x_3)$.

Confusingly, the literature sometimes refers to $\mathcal{T}$-models as $\mathcal{T}$-algebras [Mac Lane 2010, §V.6]. To avoid this confusion, we will use the terminology '$\mathcal{T}$-models' as much as possible, with one exception. When $\mathcal{T}$ is implicit, we will sometimes use the portmanteau 'fral' of '**fr**ee **al**gebra'. We prefer it over alternatives ('frmo' or 'frem').

In summary, fral simplification provides a guiding principle for designing the data structure in the middle of (1): it is an implementation of the fral. The success criterion for the design of the simplification code is the universal property, which singles the fral up to a unique isomorphism of models over X. The different components of the universal property provide the following methodology for this creative process. Failure in each step in this process may provide insight into how the current data structure fails, and how it may be revised:

(1) *Equip the data structure with a $\Sigma_\mathcal{T}$-algebra structure.* Failure in this step typically exposes inputs this simplifier cannot deal with. Revise the data structure to support these operations.

(2) *Prove it validates the axioms, i.e., forms a $\mathcal{T}$-model.* Failure in this step exposes equations the structure does not use, i.e., completeness issues. Revise the data structure to use the additional equations, or relax the presentation, giving up on using these equations.

(3) *Define a function out of it to any other $\mathcal{T}$-model over* x. Failure in this step typically exposes 'junk' in the data structure: components that are not captured by the presentation such as missing operators. Revise the data structure to exclude the junk, or extend the presentation with more operators to account for the additional components.

(4) *Show this function is a homomorphism.* Failure in this step typically exposes unprovable equations that the data structure validates. Revise the data structure/presentation.

(5) Show this homomorphism is unique. Failure in this step similarly exposes implicit axioms in the data structure, revise data structure/presentation.

While each step requires mathematical creativity, the overall structure breaks the task down into smaller steps, a checklist. This checklist organises the simplification code along these steps, making simplifier development more methodical.

We now turn to frex simplification, through the second example from page 2:

$$-6 + (x + 3) + (y + x) \xmapsto{\text{evaluate}_\mathbb{Z}} (-3, 2, 1) \xmapsto{\text{reify}} -3 + 2x + y \tag{2}$$

Let A be a $\mathcal{T}$-model and X a set whose elements we treat as representing variables. An *extension* of A by X is a triple a = (a.Model, a.Var, a.Embed.H) consisting of a $\mathcal{T}$-model a.Model, a function a.Var : X ⟶ U(a.Model), and a $\mathcal{T}$-homomorphism a.Embed : A ⟶ a.Model. This concept lets us phrase the distinction between meta-level constants such as 3 and 6, and meta-level variables such as y in the meta-level expression $-6 + (x + 3) + (y + x)$ from (2). Taking A := $\mathbb{Z}$ and X := $\{x, y\}$, both the concrete elements of the algebra A (e.g., 3 and -6) and the abstract variables from X are part of the vocabulary of the simplification process. It uses the evaluation equation $-6 + 3 = -3$.

A *morphism* h : a ⟶ b of extensions of A by X is a $\mathcal{T}$-homomorphisms h : a.Model ⟶ b.Model that makes the two triangles in the diagram on the right commute. A *free extension* of A by X is then an extension from which there is a unique such morphism to every other extension of A by X. For example, the free extension of a commutative monoid A by two variables X := $\{x, y\}$ is the $\Sigma_{(2,0)}$-algebra over U A × $\mathbb{N}^2$, given componentwise, equipped with the $\Sigma_{(2,0)}$-homomorphism that sends $u \in$ U A to $(u, 0, 0)$ and the function that sends x to $(A\llbracket 0 \rrbracket, 1, 0)$ and y to $(A\llbracket 0 \rrbracket, 0, 1)$. The unique morphism to any extension (B, e, b) sends (u, a_0, a_1) to $bu + a_0(e\,x) + a_1(e\,y)$.

Figure 1 summarises the frals and frexes in this article. All but the last frex are well-known representations. Our contribution is to implement, alongside these representations, constructive proofs that they represent the fral or the frex. These proofs require the universal algebraic concepts we introduced so far. Other simplification modules can reuse these concepts and the representation theorems for existing simplifiers. This design enables extensible and modular user and library code.

2.3 Setoids

Designing Frex around the concepts from Section 2.1 and Section 2.2 supports term simplification in a uniform way. Generalising the design further to use *setoids* [Bishop 1967; Hofmann 1997] instead of sets enables the same abstractions to offer more flexible functionality: printing terms and proofs, proof simplification, code generation, and more, which Sections 5 and 6 explore. The ability to extract such intensional information from terms violates algebraic equality, and differentiates this design from work involving quotients, for example quotient inductive-inductive types (QIITs) [Altenkirch

	monoids	*commutative* **monoids**	*involutive* **monoids**
free algebra	lists of variables $$yxxyx$$	origin-intercepting linear polynomials $$a_1x_1+\ldots+a_nx_n$$ $(a_i : \mathtt{Nat})$	lists over ordinary & singly-involuted variables $$y\overline{x}xx\overline{y}x$$
free extension	alternating lists in $\mathbb{M}_{2\times2}(\mathtt{Nat})[y]$ $$\begin{pmatrix}1 & 3\\0 & 2\end{pmatrix}y\begin{pmatrix}0 & 1\\1 & 0\end{pmatrix}y$$	linear polynomials in $A[x_1,\ldots,x_n]$ $$c+a_1x_1+\ldots+a_nx_n$$ $(a_i : \mathtt{Nat}, c : A)$	alternating lists with tagged variables in $\mathtt{String}[x,y]$ $$\texttt{""}x\texttt{"hello"}y\texttt{"olleh"}\overline{x}\texttt{""}$$

Fig. 1. Frals and frexes for varieties of monoids

$$\frac{}{\mathsf{X} \vdash t = t}\text{REFL} \qquad \frac{\mathsf{X} \vdash s = t}{\mathsf{X} \vdash t = s}\text{SYM} \qquad \frac{\mathsf{X} \vdash t = s \quad \mathsf{X} \vdash s = r}{\mathsf{X} \vdash t = r}\text{TRANS} \qquad \frac{(\mathsf{X} \vdash t = s) \in \mathcal{T}.\mathtt{Axiom}}{\mathsf{X} \vdash t = s}\text{AX}$$

$$\frac{\mathsf{Y} \vdash t = s \quad \theta_1, \theta_2 : \mathsf{Y} \longrightarrow \mathtt{Term}\ \mathsf{Y} \quad (\mathsf{X} \vdash \theta_1 y = \theta_2 y)_{y \in \mathsf{Y}}}{\mathsf{X} \vdash t[\theta_1] = s[\theta_2]}\text{CONG} \qquad \frac{\mathsf{UA} \vdash t}{\mathsf{X} \vdash t = A[\![t]\!]}\text{EVAL}$$

Fig. 2. Provability in (a) equational logic (unshaded) and (b) the evaluation rule

et al. 2018; Altenkirch and Kaposi 2016; Kaposi et al. 2019]. This section presents some mathematical background about setoids and their realisation in Idris2.

A *setoid* $\mathsf{X} = (\mathsf{U}\,\mathsf{X}, (\equiv_{\mathsf{X}}))$ consists of a set $\mathsf{U}\,\mathsf{X}$ and an equivalence relation $(\equiv_{\mathsf{X}})$ over $\mathsf{U}\,\mathsf{X}$. A *setoid homomorphism* $f : \mathsf{X} \leadsto \mathsf{Y}$ is a relation-preserving function between sets X and Y: $f\,x \equiv_{\mathsf{Y}} f\,y$ whenever $x \equiv_{\mathsf{X}} y$. We think of elements in $\mathsf{U}\,\mathsf{X}$ as representatives of the equivalence classes of $(\equiv_{\mathsf{X}})$, and so every setoid homomorphism induces a (unique) function between the quotients $\mathsf{X}/(\equiv_{\mathsf{X}}) \to \mathsf{Y}/(\equiv_{\mathsf{Y}})$. One way to construct a setoid is to equip a set X with its equality relation to give $(\mathsf{X}, (=))$, but using setoids also allows us to explicate sophisticated equivalence relations and define operations on them that are not supported by the corresponding quotients. For example, given a presentation $\mathcal{T}$ and a set X, we can define the *provability* relation $\mathsf{X} \vdash - = -$ over terms $\mathsf{X} \vdash t$ (cf. Fig. 2). Related elements in the setoid $(\mathtt{Term}\,\mathsf{X}, \mathsf{X} \vdash - = -)$ represent different, but provably equal, terms. In contrast, elements in the quotient $\mathtt{Term}\,\mathsf{X}/(\mathsf{X} \vdash - = -)$ represent equivalence classes of provably equal terms. Figure 2 also includes the evaluation axiom (EVAL), which we will use to present the frex.

Setoids and their homomorphisms form a common technique to complete an intensional type theory with extensional functions and quotients, requiring users to establish that every defined function is a setoid homomorphism. However, in FREX we make essential use of setoids that quotient types [Hofmann 1995] do not allow. For example, the terms-up-to-provability setoid supports operations such as $\mathtt{vars} : \mathtt{Term}\,\mathsf{X} \to \mathtt{List}\,\mathsf{X}$ which extract the list of variables appearing in a given term in-order. This function is not a setoid homomorphism. The quotient can only support such extraction following a canonicalisation. Both flavours of function are useful in applications, but setoids, and not quotients, support both.

In FREX, we implement universal algebra *internal* to setoids: carriers are setoids; algebraic operations are setoid homomorphisms; algebra homomorphisms and environments must also be

```
record Equivalence (A : Type) where
  constructor MkEquivalence
  0 relation: Rel A
  reflexive : (x        : A) -> relation x x
  symmetric : (x, y     : A) -> relation x y -> relation y x
  transitive: (x, y, z : A) -> relation x y -> relation y z
                            -> relation x z

record Setoid where
  constructor MkSetoid
  0 U : Type
  equivalence : Equivalence U
```

```
data Setoid : Type where
  MkSetoid : (0 U : Type) ->
    (equivalence
      : Equivalence U) -> Setoid

0
U : Setoid -> Type
U (MkSetoid x _) = x

equivalence : (s : Setoid) ->
            Equivalence (U s)
equivalence (MkSetoid _ y) = y
```

Fig. 3. (a) Equivalence relations and setoids as records and (b) example desugaring into a GADT and projections

setoid homomorphisms; and the unique homomorphisms in the universal properties of the fral and the frex must be setoid homomorphisms. With this generalisation, the setoid $\mathrm{Term}\ X/(X \vdash - = -)$ also satisfies the fral universal property constructively. Therefore, there is a unique canonical setoid isomorphism to every other fral. Similarly, by including the evaluation equations (EVAL), we obtain a setoid frex together with a setoid isomorphism to every other frex. These isomorphisms let us extract simplification proofs generically out of user-defined simplifiers. We discuss the decision to use setoids in Section 10 where we refer to the concrete benefits setoids offer.

Setoids in Idris2. We represent equivalence relations and setoids in Idris2 with *records* in Fig. 3a. Idris2 records are syntactic sugar for a single-constructor data declaration and automatically generated *field projections*, as in Fig. 3b. Idris2 also automatically generates the post-fix projections for each field using a dotted notation, writing b.equivalence.relation for the nested projection. The annotation 0 preceding the definition of the field U is a *quantity* [Atkey 2018; McBride 2016] indicating that the field is not represented at runtime, but may be used in types. Readers can safely ignore these annotations. There is only a handful of them in this article. We maintain them to demonstrate and emphasise that Frex does not require us to retain, for example, runtime representations of types. If you are reading this article in colour, our listings include semantic highlighting, designating the semantic class of each lexeme: **data** constructor, type constructor, defined function or value, and variable in a binding/bound occurence. We define setoid homomorphisms:

```
SetoidHomomorphism : (a,b : Setoid)
  -> (f : U a -> U b) -> Type
SetoidHomomorphism a b f
  = (x,y : U a) -> a.equivalence.relation x y
  -> b.equivalence.relation (f x) (f y)
```

```
record (~>) (A,B : Setoid) where
  constructor MkSetoidHomomorphism
  H : U A -> U B
  homomorphic : SetoidHomomorphism A B H
```

The Appendix includes expanded examples for setoids of functions and quotient setoids.

This technique is affectionately dubbed 'setoid hell', since we need to prove that all our functions are setoid homomorphisms. Following Hu and Carette [2021], we manage setoid hell by structuring code categorically, organising results into homomorphisms between appropriate setoids.

3 Universal Algebra in Frex

To define an interface to algebraic simplifiers, we first specify and represent algebraic structures. We implement signatures and their operators in Frex as follows (below, left and middle):

```
(^) : Type -> Nat -> Type
(^) a n = Vect n a
algebraOver : (sig : Signature)
  -> (a : Type) -> Type
sig `algebraOver` a =
  (f : Op sig) -> a ^ (arity f) -> a
record Algebra (Sig : Signature) where
  constructor MakeAlgebra
  0 U   : Type
  Semantics : Sig `algebraOver` U
```

```
CongruenceWRT : {n : Nat} -> (a : Setoid) ->
  (f : (U a) ^ n -> U a) -> Type
CongruenceWRT a f = SetoidHomomorphism (VectSetoid n a) a f
record SetoidAlgebra (Sig : Signature) where
  constructor MkSetoidAlgebra
  algebra : Algebra Sig
  equivalence : Equivalence (U algebra)
  congruence : (f : Op Sig) ->
    (MkSetoid (U algebra) equivalence)
      `CongruenceWRT` (algebra.Sem f)
```

Fig. 4. Algebras and setoid algebras in FREX

```
record Signature where
  constructor MkSignature
  OpWithArity : Nat -> Type
```

```
record Op (sig : Signature) where
  constructor MkOp
  {arity : Nat}
  snd : sig.OpWithArity arity
```

```
data Operation
     : Nat -> Type where
  Neutral : Operation 0
  Product : Operation 2
```

The implementation uses Idris2's implicit record field for `arity`. Users define concrete instances of `Signature`, such as the signature `MkSignature Operation` for monoids, by defining a type family for the indexed field `OpWithArity` (above, right).

FREX represents the domain of an *n*-ary operation with an *n*-ary vector (Fig. 4). (As in Haskell, backticks turn any name into an infix operator.) For example, the additive natural numbers form an algebra for the monoid signature as follows:

```
Additive : Algebra Monoid.Theory.Signature
Additive =  MkAlgebra {U = Nat, Sem = \case
  Neutral => 0
  Product => plus}
```

Instead of specifying the `Semantics` field directly, this code uses the smart constructor `MkAlgebra`, which has a `Sem` argument, instead of `Semantics`. This smart constructor transfers its `Sem` argument into `MakeAlgebra`'s `Semantics` field by uncurrying each `n`-ary function into a function taking an `n`-ary vector of arguments. The departure from our usual naming scheme in which it is the constructor of a record `R` that is called `MkR` indicates that the smart constructor `MkAlgebra`, not the record constructor `MakeAlgebra`, is the preferred way to construct `Algebra` instances. The `\case` keyword is an anonymous function that immediately pattern-matches its argument. *Setoid algebras* further require an equivalence relation that forms a congruence w.r.t. the operations (Fig. 4).

We implement terms over a signature as follows, mirroring their mathematical definition:

```
data Term : (0 sig : Signature) -> Type -> Type where
  ||| A variable with the given index
  Done : {0 sig : Signature} -> a -> Term sig a
  ||| An operator, applied to a vector of sub-terms
  Call : {0 sig : Signature} -> (f : Op sig) ->
         Vect (arity f) (Term sig a) -> Term sig a
```

Terms form an algebra, the *free* algebra, with symbols denoting term formers:

```
data Axiom = LftNeutrality
           | RgtNeutrality
           | Associativity
MonoidTheory : Presentation
MonoidTheory = MkPresentation Theory.Signature Theory.Axiom $ \case
                 LftNeutrality => lftNeutrality Neutral Product
                 RgtNeutrality => rgtNeutrality Neutral Product
                 Associativity => associativity        Product
MonoidStructure : Type
MonoidStructure = SetoidAlgebra Signature
Monoid : Type
Monoid = Model MonoidTheory
```

Fig. 5. Axiomatising monoids in Frex

```
Free : (0 sig : Signature) -> (0 x : Type) -> Algebra sig
Free sig x = MakeAlgebra (Term sig x) Call
```

Terms also form a monad, with `Done` as its unit and substitution as its sequencing operation.

Turning to equations, Frex only needs equations in a finite context, and we call its cardinality the *support* of the equation. We implement equations and presentations as follows:

```
record Equation              record Presentation where     associativity : {sig : Signature}    1
  (Sig : Signature) where      constructor MkPresentation     -> EqSpec sig [2]                    2
  constructor MkEq             signature : Signature        associativity product =              3
  support  : Nat              0 Axiom   : Type               let (+) = call product in            4
  lhs, rhs :                   axiom     : (ax : Axiom) ->   MkEquation 3 $ X 0 + (X 1 +  X 2)    5
    Term Sig (Fin support)      Equation signature                   =-= (X 0 +  X 1) + X 2      6
```

For example, the monoid presentation **Monoid** has three axioms: left and right neutrality, and associativity. Frex defines a generic collection of axiom schemes (above, right). The second argument to `EqSpec` lists the arities `[n1, ..., nk]` of the operations in a scheme, so on lines 1–2 the argument `[2]` in the type `EqSpec sig [2]` states that the scheme involves a single binary operation. The first argument `3` in `MkEquation 3` (lines 5–6) indicates that the declaration of the scheme involves three variables `X 0`, `X 1`, and `X 2`. We use these generic axiom schemes to construct, for example, the theory of monoids in the middle of Fig. 5.

Fig. 6 shows Frex's representation of what it means for an algebra to validate an equation. We use Idris2's dependent pairing construct to pair an algebra with an environment in the standard entailment syntax . The following code validates the monoid axioms for our running example:

```
IsMonoid : Validates MonoidTheory NatAdditive
IsMonoid LftNeutrality env = Refl
IsMonoid RgtNeutrality env = plusZeroRightNeutral _
IsMonoid Associativity env = plusAssociative _ _ _
```

We define models for a presentation:

```
models : {sig : Signature} -> (a : SetoidAlgebra sig) -> (eq : Equation sig) ->
         (env : Fin eq.support -> U a.algebra) -> Type
models a eq env = a.equivalence.relation (a.Sem eq.lhs env) (a.Sem eq.rhs env)

(=|) : {sig : Signature} -> (eq : Equation sig) ->
       (a : SetoidAlgebra sig ** Fin eq.support -> U a.algebra) -> Type
eq =| (a ** env) = models a eq env

ValidatesEquation : (eq : Equation sig) -> (a : SetoidAlgebra sig) -> Type
ValidatesEquation eq a = (env : Fin eq.support -> U a.algebra) -> eq =| (a ** env)

Validates : (pres : Presentation) -> (a : SetoidAlgebra pres.signature) -> Type
Validates pres a = (ax : pres.Axiom) -> ValidatesEquation (pres.axiom ax) a
```

Fig. 6. Equational validity in an algebra

```
record Model (Pres : Presentation) where
  constructor MkModel
  Algebra  : SetoidAlgebra (Pres).signature
  Validate : Validates Pres Algebra
```

We can now define a monoid to be a **Monoid**-model, as in Fig. 5. For another example, now putting everything together, we validate the monoid structure of multiplication as follows:

```
Multiplicative : Monoid                                        1
Multiplicative = MkModel                                       2
  { Algebra = cast {from = Algebra Signature} $                3
      MkAlgebra {U = Nat, Sem = \case Neutral => 1             4
                                      Product => mult}          5
  , Validate = \case                                           6
      LftNeutrality => \env => plusZeroRightNeutral _          7
      RgtNeutrality => \env => multOneRightNeutral _           8
      Associativity => \env => multAssociative _ _ _           9
  }                                                            10
```

The coercion function (`cast : from -> to`) on Line 3 is a method in the standard Idris type-class/interface `Cast from to`. FREX exports a `Cast`-instance that converts an algebra into a setoid algebra whose equivalence relation is propositional equality (=). Lines 7–9 use results about the natural numbers from Idris2's standard library.

Using FREX. Unless they are already working abstractly with an algebraic structure, we expect that in practice users start by recognising that their concrete algebra validates the axioms of an existing simplification module—*frexlet* for short. Such modules export a presentation, convenient notation suites for its signature, and fral and/or frex simplifiers for this presentation.

As a concrete example, we will take computations with lists that also involve the `reverse` function. These form an *involutive* monoid: a monoid A equipped with a unary *involution* operator $x \mapsto \overline{x}$: $UA \rightarrow UA$ satisfying two axioms $\overline{\overline{x}} = x$ and $\overline{xy} = \overline{y}\,\overline{x}$. We then equip our type of interest, lists, with an involutive structure as in Fig. 7. We can use this algebra and the involutive monoid to discharge equations containing list variables and concrete lists:

```
ListInvMonoid : {0 a : Type} -> InvolutiveMonoid
ListInvMonoid = MkModel
  { Algebra = cast $ MkAlgebra
      {sig = Monoid.Involutive.Theory.Signature}
      { U = List a                                        -- Carrier
      , Sem = \case                                       -- Operations
          Mono monoidOp => case monoidOp of               --   Inherited from monoids
              Neutral => []
              Product => (++)
          Involution => reverse
      }
  , Validate = \case                                      -- Validate equations
      Mon LftNeutrality  => \env => Refl                  -- Directly, or
      Mon RgtNeutrality  => \env => appendNilRightNeutral _  -- use existing standard
      Mon Associativity  => \env => appendAssociative _ _ _  -- library functions
      Involutivity       => \env => reverseInvolutive _
      Antidistributivity => \env => sym (revAppend _ _)
  }
```

Fig. 7. The involutive monoids of list reversal

```
lemma : {x : List a} -> (i,j : a) ->                                              1
    reverse ([i] ++ reverse x ++ []) ++ [j]                                       2
    =                                                                             3
    x ++ [i, j]                                                                   4
lemma i j = solve 1 (Involutive.Frex.Frex ListInvMonoid) $                        5
    ((Sta [i] .*. (Dyn 0) .inv .*. I1) .inv) .*. Sta [j]                          6
    =-=                                                                           7
    Dyn 0 .*. Sta [i, j]                                                          8
```

The solve function takes as argument the number of variables (n=2 on line 2) in the algebraic term to simplify, and an algebraic simplifier from the frexlet (Involutive.Frex.Frex on line 4). The final argument is a pair of terms with n=2 variables (Dyn 0 and Dyn 1) and concrete values from the algebra. By importing notation modules the frexlet provides, we can use infix multiplicative notation such as (.*.). The type-checker then infers the terms to substitute for each variable.

In this example, we used solve to define a stand-alone lemma, but we may also call solve directly from a chain of equational reasoning steps. When we extract lemmas, we often want to prove them more abstractly, for *all* involutive monoids. In that case we use a fral:

```
ExampleFral : {a : InvolutiveMonoid} -> (x,y,z : U a)                             1
  -> let %hint notation : ?                    -- Open notation hints for the monoid   2
         notation = a.Notation1                -- for infix operator (.*.) and         3
     in a.rel                                  -- postfix operator (.inv)              4
       (x .*. y.inv .*. z).inv                                                         5
       (z.inv .*. y .*. x.inv)                                                         6
ExampleFral x y z =                                                               7
  let %hint notation : ?                       -- ditto, but for terms                8
      notation = Involutive.Notation.multiplicative1                              9
```

```
in solve 3 (Involutive.Free.FreeInvolutiveMonoidOver 3) $              10
(X 0 .*. (X 1).inv .*. X 2).inv =-= (X 2).inv .*. X 1 .*. (X 0).inv    11
```

Lines 2–3 and 8–9 overload the infix and postfix notation using the frexlet's built-in notation suites. Concretely, the projection `Notation1` brings into scope the functions `(.*.)` and `(.inv)` when writing algebraic terms. The `solve` function takes the number of free variables and a corresponding fral simplifier (line 10), as well as the two terms representing the equation of interest. The variables `x`, `y`, `z` (bound in line 7) are implicitly used in this call. §4.2 covers the type of `solve` in more detail.

4 Free Extensions and Free Models/Algebras

Before delving into the details of FREX's core, we revisit our frexlet representations using examples for elements in the fral and the frex for ordinary, commutative, and involutive monoids (see Fig. 1).

The elements in the free monoid are lists of the variables appearing in the term. The elements in the free extension of a monoid are lists alternating between concrete elements in the given monoid, and freely adjoined variables. The figure shows an element in the free extension by 1 variable (y) of the multiplicative monoid of 2×2 matrices with natural-number components. The matrix y is unknown, and so its occurrence separates the elements in the list.

Further assuming commutativity equates more terms, resulting in a representation of the free commutative monoid over `n` variables as an `n`-vector of coefficients, representing a linear polynomial. Freely extending a commutative monoid `A` by `n` variables can be represented by a concrete coefficient c : `A` together with an `n`-vector of coefficients, representing a linear polynomial over `A`.

If we instead include an involutive operation $x \mapsto \overline{x}$ over the monoid, we get lists of variables and alternating lists whose letters may be tagged as involuted. The figure demonstrates the free extension of the monoid structure of `String` concatenation, with string reversal for the involution.

4.1 Universal Properties

We now show the Idris2 realisation of the universal algebraic concepts from Section 2. In the previous Section 3, we introduced these definitions:

- An `Algebra` for a signature, which consists of a carrier and an interpretation of operations (Fig. 4). It does not have to satisfy any equation.
- A `SetoidAlgebra` is simultaneously a setoid structure (an equivalence relation) and an algebra structure over the same carrier, whose operations are setoid homomorphisms (Fig. 4).
- A (setoid) `Model` for a presentation is a setoid algebra validating the presentation's axioms.

In this section, we introduce these definitions:

- Homomorphisms of setoid algebras.
- A `ModelOver` a setoid `X` is a model `A` equipped with a setoid homomorphism `X` to `A`.
- An `Extension` of a model `A` by a setoid `X` is another model `B` together with a setoid homomorphism from `X` to `B` and a model homomorphism from `A` to `B`.
- Their morphisms.

We will then show how we define in Idris2 both the free model (fral) over a setoid `X` and the free extension (frex) of `A` by `X` as the initial ones.

FREX defines homomorphisms of setoid algebras in Fig. 8, by requiring the underlying function to be a setoid homomorphism between the corresponding setoids. The code uses an appropriate `cast` function (see page 11) that assembles these setoids from the data in each setoid algebra. Each `a` : `Algebra` sig defines a homomorphic extension operator `a.Sem` : `Term sig x -> (x -> U a) -> U a` by structural induction over the term (i.e. folding). The following term valuates to `5+0+7`, e.g.:

$$(\texttt{Nat.Additive}).\texttt{Sem} \ (\texttt{X 0.+.O1.+.X 1}) \ (\backslash\texttt{case} \ \{\texttt{0=>5; 1=>7}\})$$

```
Preserves : {sig : Signature}            Homomorphism : {sig : Signature}
   -> (a, b : SetoidAlgebra sig)            -> (a, b : SetoidAlgebra sig) -> (h : U a -> U b) -> Type
   -> (h : U a -> U b)                    Homomorphism a b h = (f : Op sig) -> Preserves a b h f
   -> (f : Op sig) -> Type
Preserves {sig} a b h f
   = (xs : Vect (arity f) (U a))
      -> b.equivalence.relation        record (~>) {Sig : Signature} (a, b : SetoidAlgebra Sig) where
            (h $ a.Sem f xs)              constructor MkSetoidHomomorphism
            (b.Sem f (map h xs))         H : cast {to = Setoid} a ~> cast b
                                         preserves : Homomorphism a b (.H H)
```

Fig. 8. Setoid algebra homomorphisms in Frex

```
record ModelOver                         record Extension {Pres : Presentation}
    (Pres : Presentation)                    (A : Model Pres)(X : Setoid) where
    (X : Setoid) where                     constructor MkExtension
  constructor MkModelOver                  Model : Model Pres
  Model : Model Pres                       Embed : A ~> Model
  Env : X ~> cast Model                    Var   : X ~> cast Model
PreservesEnv : {Pres : Presentation}     record (~>) {Pres : Presentation}
    -> {X : Setoid}                          {A : Model Pres} {X : Setoid}
    -> (a, b : Pres `ModelOver` X) ->         (Extension1, Extension2 : Extension A X) where
    (cast {to = Setoid} a.Model            constructor MkExtensionMorphism
      ~> cast b.Model) -> Type             H : (Extension1).Model ~> (Extension2).Model
PreservesEnv a b h =                       PreserveEmbed :
 (X ~~> cast b.Model).equivalence.relation    (cast A ~~> (Extension2).Model)
    (h . a.Env) b.Env                          .equivalence.relation
record (~>)                                       (H . (Extension1).Embed)
    {Pres : Presentation} {X : Setoid}                (Extension2).Embed
    (A, B : Pres `ModelOver` X) where      PreserveVar   :
  constructor MkHomomorphism                  (X ~~> cast (Extension2).Model)
  H : (A .Model) ~> (B .Model)                 .equivalence.relation
  preserves : PreservesEnv A B (H .H)             ((H).H . (Extension1).Var)
                                                        (Extension2).Var
```

Fig. 9. Structure and its preservation for (a) models over a setoid, and (b) extensions of a model

Similarly, Fig. 9 presents the declarations for models over a setoid and extension of a model by a setoid. It expresses the equations in the commuting diagrams through the extensionality relation on the setoid of functions from Fig. 15b in Appendix B and the power of a model by a setoid (see §4.3). As Idris2 supports type-directed disambiguation, we overload the record name (~>).

The *free* model over a set (fral) and the *free* extension (frex) of an algebra by a set is then the initial such structure: there is a unique structure-preserving map from the free structure to every structure. This succinct definition, while standard, packs much structure. By way of introduction, we will unpack it for the free commutative monoid over Fin n, the finite set with n elements.

First, we designate a commutative monoid for the the fral. This structure is the data structure our simplifier will use to represent the equivalence classes of terms. In Fig. 1, we mentioned the

carrier consists of origin-intercepting linear polynomials with `Nat` coefficients $p = a_1 x_1 + \ldots + a_n x_n$, which we represent with `n`-tuples of natural numbers and pointwise addition:

```
Carrier : (n : Nat) -> Setoid
Carrier n = VectSetoid n
                (cast Nat)
```

$$0 := 0x_1 + \ldots + 0x_n \qquad p{+}q := (a_1 + b_1)x_1 + \ldots + (a_n + b_n)x_n$$
$$:= \texttt{[0,\ldots,0]} \qquad\qquad := \texttt{[}a_1{+}b_2, \ldots, a_n{+}b_n\texttt{]}$$
$$= \texttt{replicate n 0} \qquad\qquad = \texttt{map (uncurry (+)) (zip as bs)}$$

Denote the resulting `CommutativeMonoid` by `Model n`. For the `Env` component, use tabulation to define `unit n : Fin n → Carrier n`, with `1` in the argument position and `0` elsewhere:

$$\texttt{unit n } i := 1x_i \qquad\qquad\qquad \text{where}$$
$$:= \texttt{[0,\ldots,0,1,0,\ldots,0]} \qquad \texttt{kronecker i j} := \begin{cases} i = j : & 1 \\ i \neq j : & 0 \end{cases}$$
$$= \texttt{tabulate \$ kronecker i}$$

The initiality of this structure follows from the normal form property — every origin-intersecting linear polynomial p can be represented as $p = \sum_{i=1}^{n} a_i \cdot \texttt{unit n i}$:

```
normalForm : (n : Nat) -> (xs : U (Model n)) ->
    xs = (Model n).sum (tabulate $ \i => (index i xs) *. (unit n i))
```

Since monoid homomorphisms preserve the summation and multiplication-by-a-natural, the unique structure preserving map `h : (Model n, unit n) → a` is this homomorphism:

```
h xs = a.Model.sum (mapWithPos (\i,k => k *. a.Env.H i) xs)
```

This standard argument lies behind many simplifiers, as well as more advanced techniques like normalisation-by-evaluation. Frex takes the same approach, but also explores how to use general-purpose constructions involving frals and frexes, and bespoke facts about algebraic structures, to construct new frals and frexes.

To summarise, to implement a fral/frex simplifier, the developer follows these steps:

- Design a data-structure for the carrier of the frex/fral's algebra, e.g. for commutative monoids: `Vect n Nat` for the fral and (`U a`, `Vect n Nat`) for the frex.
- Equip it with a setoid algebra structure: pointwise operations with propositional equality.
- Equip it with the appropriate additional structure, e.g. the unit for the fral and the `Var`iable function and the `Embed`ding homomorphism for the frex.
- Define the function underlying the homomorphism into any other algebra over the variable setoid or extension, e.g. linear combination for commutative monoids.
- Prove that this function is a homomorphism and its uniqueness.

These steps realise the simplifier development methodology we described in Section 2.2 (page 5).

4.2 Solver Interface

We can now explain the interface Frex gives to the `solve` functions. We describe the frex-based interface in detail; the fral-based interface is similar. We implement the core functionality in the auxiliary function `solveVect` in Fig. 10.

The argument `frex` (line 2) is an implementation of a frex simplifier for some `pres`-model `a`, extended with `n` free variables (line 1). We erase the number of variables at runtime, and so we also erase the type `Fin n`. We cast an erased type to a setoid instead of an unerased type, i.e.:

```
irrelevantCast : (0 a : Type) -> Setoid    instead of    cast : (a : Type) -> Setoid
```

The function also takes an `env`ironment of terms to substitute for the free variables in the simplification equation (line 2). In this auxiliary function, we present the environment using an `n`-ary

```
1   solveVect : {0 n : Nat} -> {pres : Presentation} -> {a : Model pres} ->
2     (frex : Frex a (irrelevantCast $ Fin n)) -> (env : Vect n (U a)) ->
3     (eq : ( Term pres.signature (U a `Either` Fin n)
4           , Term pres.signature (U a `Either` Fin n))) ->
5     {auto prf : frex.Data.Model.rel
6         (frex.Sem (fst eq) (frexEnv {x = cast $ Fin n} frex).H)
7         (frex.Sem (snd eq) (frexEnv {x = cast $ Fin n} frex).H)}
8         ->
9     a.rel (a.Sem (fst eq) (either Prelude.id (flip Vect.index env)))
10          (a.Sem (snd eq) (either Prelude.id (flip Vect.index env)))
```

Fig. 10. Core frex-based simplification routine

vector of terms over the algebra's carrier. Next comes the `eq`uation we want to discharge (line 3), involving either concrete values (of type `U a`) and any of the `n` available variables. Both the `frex` and the `a`lgebra with its `env`ironment give rise to extensions in the formal sense, which we can use to give an environment for the `eq`uation in question, namely a setoid homomorphism from the joint setoid of constants and free variables to the carrier of the model underlying the extension:

```
extEnv : {a : Model pres} -> {x : Setoid} -> (ext : Extension a x) ->
  Either (cast a) x ~> cast ext.Model
extEnv ext = either ext.Embed.H ext.Var
```

where: `either : {0 a, b, c : Setoid} -> (a ~> c) -> (b ~> c) -> (a `Either` b) ~> c`

We use these environments to interpret the equation, once in the `frex` (lines 6–7) and once in the given algebra (lines 9–10). If the equation holds in all extensions, it will hold in the `frex` and in `a`, and, moreover, homomorphisms of extensions will preserve this interpretation. Interpreting this equation in the `frex` may have better decidability properties over equivalence in `a`.

We use Idris2's auto-implicits mechanism to search for the equivalence of the `frex` interpretations. This mechanism will try to find terms that resolve the implicit argument `prf`, using a heuristic informed by unification, that will also attempt to apply data constructors.

Typically, Idris2's judgemental equality decides the `frex` setoid's equivalence relation, and the number of variables we extend by is known statically. This case can happen when the equivalence relation on the setoid algebra `a` is decidable by judgemental equality. Then, the type of the `prf` argument (line 5) is a propositional equality between closed terms. Judgemental equality decides this relation between the interpretations in the type of `prf`. In Idris2, the auto-search heuristic tries to use **Refl**, and promotes the required equation to a judgemental equality constraint. Even when the setoid relation is not decidable by judgemental equality, making `prf` an auto-implicit may provide more functionality in the future. We may be able to freely extend algebras whose propositional equality is only partially decidable by judgemental equality (e.g. function types in a type theory with function extensionality), or by a sophisticated decision procedure (e.g. multiset equality).

We use metaprogramming abstractions (Fig. 11) to simplify the user-facing interface. The `PI` combinator produces an `n`-ary telescope of **Visible**/**Hidden**/**Auto** arguments, packaged as an `n`-ary vector which it passes this its argument `b`. Using this abstraction to reduce `solve` (Fig. 12) to `solveVect` (Fig. 10). Compare their types: `solve` curries the `env`ironment `n`-vector into `n` implicit arguments. Unification can resolve these arguments to the free variables in the discharged equation.

```
data Visibility = Visible | Hidden | Auto
Pi : Visibility -> (a : Type) -> (a -> Type) -> Type
Pi Visible a b = (x : a) -> b x
Pi Hidden  a b = {x : a} -> b x
Pi Auto    a b = {auto x : a} -> b x
PI : (n : Nat) -> Visibility -> (a : Type) -> (Vect n a -> Type) -> Type
PI Z     vis a b = b []
PI (S n) vis a b = Pi vis a (\ x => PI n vis a (b . (x ::)))
```

Fig. 11. Metaprogramming abstractions for curried Π-types

```
1   solve : (n : Nat) -> {pres : Presentation} -> {a : Model pres} ->
2       (frex : Frex a (cast $ Fin n)) ->
3       PI n Hidden (U a) $ (\ env =>
4       (eq : ( Term pres.signature (U a `Either` Fin n)
5             , Term pres.signature (U a `Either` Fin n))) ->
6       {auto prf : frex.Data.Model.rel
7          (frex.Sem (fst eq) (frexEnv frex).H)
8          (frex.Sem (snd eq) (frexEnv frex).H)}
9          ->
10      a.rel (a.Sem (fst eq) (either Prelude.id (flip Vect.index env)))
11            (a.Sem (snd eq) (either Prelude.id (flip Vect.index env))))
```

Fig. 12. User-facing frex-based simplification routine

4.3 Powers

The commutative monoid structure `Model n` instantiates a general construction: $\mathcal{T}$-models have powers by setoids. The *power* of an algebra `A` by a set(oid) `X` is the terminal *parameterisation*. Parameterisations, shown succinctly on the right, are an `X`-indexed collection of algebra homomorphisms `a.Eval f : a.Model` → `A`. Requiring `a.Eval f` to be homomorphic implies that operations are given pointwise. Structure preservation uses the contravariant action `pre Hh` precomposing a homomorphism `Hh : a.Model` → `b.Model`. Universality singles out the carrier of the power as the function-space `X ~~> a.Model`. For `X = Fin n`, we can represent it by `n`-tuples from `U A`.

4.4 Frex via Coproducts with Fral

The fral and the frex are related: the free extension of `A` by `X` is the *coproduct* of `A` with the free algebra over `X`. A *cospan* `a` between A_1 and A_2 consists of two homomorphisms with a shared codomain: $A_1 \xrightarrow{\text{a.Lft}} \text{a.Sink} \xleftarrow{\text{a.Rgt}} A_2$. A *cospan homomorphism* $h : a → b$ is a homomorphism $H : \text{a.Sink} → \text{b.Sink}$ preserving the cospan as in the diagram on the right. The coproducts $A_1 \oplus A_2$ of two models A_1 and A_2 is then the initial *cospan*. All models have coproducts, but these may be difficult to represent. However, in cases such as commutative monoids, the coproduct is straightforward to represent: its carrier is the cartesian product of the components carriers.

The universal property of the frex `A[X]` combines those of the fral `Free 𝒯X` and its coproduct with `A`. Consider the following two diagrams:

The `Var`-arrows and the `Lft`-arrows correspond through the universal property of the fral: `Lft` is the unique homomorphic extension of the `Var`-arrows. The `Embed`-homomorphisms are exactly the `Rgt`-homomorphism. This identification lets us construct:

```
CoproductAlgebraWithFree pres a x : (free : Free pres x) ->
            (coprod : Coproduct a free.Data.Model) -> Frex a x
```

For commutative monoids it gives the commutative monoid of linear polynomials with natural numbers as degree-1 coefficients whose carrier is represented by (`U A`, `Vect n Nat`).

4.5 Fral via a Frex

We use the relationship between the fral and frex, described in §4.4, to derive a fral from a corresponding frex. In particular, the following calculation shows that every free algebra arises as a free extension of `Free 𝒯∅`. We use ⊎ for the disjoint union, i.e., the coproduct of sets:

$$\text{Free}\,\mathcal{T}\,X \cong \text{Free}\,\mathcal{T}\,(X \uplus \emptyset) \qquad\qquad\qquad (X \cong X \uplus \emptyset)$$
$$\cong (\text{Free}\,\mathcal{T}\,X) \oplus (\text{Free}\,\mathcal{T}\,\emptyset) \qquad (\text{Free}\,\mathcal{T}\ \text{left adjoint, left adjoints preserve coproducts})$$
$$\cong (\text{Free}\,\mathcal{T}\,\emptyset)[X]$$

Therefore, we may construct a fral from an initial algebra and its frex:

```
ByFrex : (initial : Free pres (cast Void)) ->
    Frex initial.Data.Model s -> Free pres s
```

This generic construction can produce suboptimal representations. For example, the initial monoid is easy to construct: its carrier is the unit type. Freely extending this initial monoid produces alternating lists, that interleave the unit value. Taking lists of variables instead leads to a simpler representation but requires more complicated proofs. However, generic constructions such as the foregoing allow us to trade efficiency for rapid prototyping.

4.6 Reusing Frexlets

The final example demonstrates reuse of one simplifier when constructing another. Recall the presentation of involutive monoids from the end of Section 3.

PROPOSITION 4.1 (JACOBS). *The carrier set and monoid operations of the free involutive monoid on* `X` *are those of the free monoid on the product* (`Bool`,`X`). *Similarly, the frex of an involutive monoid by* `X` *is the frex of its underlying monoid by* (`Bool`,`X`), *extended with an involution operation.*

We can prove this proposition directly, establishing the involutive axioms, and have taken this strategy in FREX. However, it is possible to prove this result without referring to the specific representation of the monoid fral/frex, resulting in a simplifier-transformer reusing the code implementing simplifiers for monoids to implement simplifiers for involutive monoid. The potential for this kind of reuse extends beyond monoids. We can phrase the result in much greater generality, and give a higher-level proof, using Jacobs's [2021] axiomatisation of involutions. This more abstract proof generalises to other notions of involutive algebras, and we plan to exploit it in the future for

generic frexlet reuse. The more abstract proof goes beyond the scope of this article, involving more abstract category theoretic notions.

5 Completeness and Certification

FREX uses setoids to extract intensional information: it automatically extracts printable representations of simplification proofs, and moreover formalises the completeness of fral/frex simplifiers. Here is how it works:

- It is easy to construct a frex and a fral using a quotient setoid: terms quotiented by provability.
- The equivalence relation in this setoid is not necessarily *effective*/*decidable*.
- An effective fral/frex must be constructed manually/creatively.
- Any fral/frex is canonically isomorphic to the quotient frex, and an effective fral/frex with a constructive universality proof has an effective isomorphism to the quotient setoid.
- In this way, a simplifier using an effective frex is constructively sound and complete.
- The universal property of the effective fral/frex lets us interpret equations it simplifies to the same value as related values in every model/extension.
- So the effective fral/frex lets us extract the data structure that represents equivalence in the setoid. For the quotient fral/frex, it is a representation of the equational proof.

Through this mechanism, there is no need to write any proof extraction code for our simplifiers. The universal properties of the fral/frex have done all the presentation-specific heavy-lifting. The remainder of this section expands this description in more technical detail (§5.1), how FREX we represent certificates (§5.2), and how we simplify proofs (§5.3).

5.1 Completeness

It is straightforward to represent the deduction rules of equational logic from Fig. 2 as a datatype `Provability` t s relating pairs of terms t and s. There are two variations on this provability relation: for free models and for free extensions.

For free models over `X`, we index the relation by $t, s \in$ `Term X`. Its definition consists of the unshaded rules in Fig. 2. The quotient setoid $Q :=$ `Term X`/`Provability` then satisfies the universal property of the free model over `X`. For the free extension of a model `a` by `X`, we index the relation by $t, s \in$ `Term` ($X \uplus \cup a$), i.e., terms with variables in the disjoint union $X \uplus \cup a$. The left-injected values represent variables from `X`. For each $c \in \cup a$, the corresponding right-injected variable, which we denote by $\underline{c}$, is a term representing c. The provability relation then adds the shaded evaluation equation from Fig. 2. It amounts to datatype constructor, for every operator $f : n$ and constants $c_1, \ldots, c_n \in \cup a$, representing the equation $f(\underline{c}_1, \ldots, \underline{c}_n) = \underline{a [\![f]\!] (c_1, \ldots, c_n)}$. The quotient setoid $Q :=$ `Term` ($X \uplus \cup a$)/`Provability` then satisfies the universal property of the frex of `a` by `X`.

The interpretation of a term t in Q is t itself, and so the quotient setoid does not help at all with simplification. Since Q satisfies the universal property, it is canonically isomorphic to every other fral/frex. So the interpretation of every two terms s, t in every fral/frex is related iff the datatype `Provability` t s is inhabited. Therefore evaluation in fral/frex is sound and complete for algebraic simplification. After designing a fral/frex S whose equivalence relation is *effective*, we can use Q as follows. The input type to the function `solve` is exactly the carrier of Q. By appealing to the universal property of the fral/frex S, we obtain an inhabitant of `Provability`, and we extracted the simplification proof, which we call the *certificate*.

5.2 Representing Certificates

FREX provides the `Lemma pres` type, consisting of a `pres.signature`-equation and derivation for it in the quotient fral for `pres`. Such lemmata are sound: every `Lemma` for a theory holds in all models of this

theory. Frex provides a `mkLemma` smart constructor which runs the given fral simplifier, constructs a proof that a stated equivalence holds, and returns a valid `Lemma`. This mechanism allows users to build up a library of lemmata for their theories. Users can then seamlessly invoke these lemmata in any model, avoiding further Frex calls.

This approach forces the user's project to depend on most of Frex indirectly through such modules. If users do not want to introduce that dependency, Frex gives them the option not to depend on it. Frex also supports proof-certificate extraction, allowing users to produce standalone libraries of ordinary Idris2 functions rather than `Lemma` independent of the Frex library. For example, the fral simplifier for monoids generates a module exporting the following Idris2 function:

```
units : (x : U m) -> O1 .+. (x .+. O1) .+. O1 =~= x
```

together with an explicit equational proof that does not call any Frex simplification steps (see Fig. 18 in the Appendix C). We now explain the certification process in technical detail.

Our goals for this extraction are to (1) produce libraries from lemmata, and (2) produce somewhat idiomatic Idris2 code. The derivation found by Frex may not be what a human would have chosen but it should nonetheless be possible for a sufficiently patient human to follow the reasoning steps.

The main challenge was to go from the rich type of derivations trees, i.e., the quotient setoid relation, to a representation that we can print in a relatively readable way. The derivation trees have arbitrarily nested transitivity, symmetry, and n-ary congruence steps. We transform them into a type of linear/flat derivations that could be pretty-printed using combinators for setoid reasoning.

We represent derivations in layers:

(a) the reflexive-transitive closure of	(b) the symmetric closure of
(c) the unary congruence closure of	(d) axiomatic reasoning steps.

We now detail each of these layers, implementing each layer by a dedicated data structure (cf. Fig. 16 in Appendix C for the full details):

(a) Reflexive-transitive closure (`RTList`): type-aligned [van der Ploeg and Kiselyov 2014] lists of steps in the closed-over relation: the target element of each list position is the source element of the next.

(b) Symmetric closure (`Symmetrise`): either the relation or its opposite.

(c) Unary congruence closure (`Locate`): It suffices to pair a term with a distinguished variable for the contextual hole, together with a step in the closed-over relation. To ease our pretty-printing code, we distinguish between using the closed-over relation in an empty context, and using it in a context with a distinguished variable represented by the Idris value `Nothing`.

(d) Axiomatic steps (`Step`): An atomic step is either a setoid equivalence, or an axiom.

Putting these together gives the type `Derivation` of linear derivations (cf. Fig. 16(e) in Appendix C).

Every derivation tree decomposes into a value in this layered representation. The modular definition of `Derivation` as a composition of the relation-transformers `RTList`, `Symmetrise`, `Locate` and `Step` makes decomposition straightforward. We use generic combinators for each closure relation-transformers. Closure under congruence is the trickiest part, decomposing an n-ary congruence in the derivation tree into n separate unary congruences, pushing them under the reflexive-transitive and symmetric closure layers, and erasing any congruence steps with the identity context.

5.3 Proof Simplication

Certification also allows us to inspect Frex-generated proofs. Frexlet developers can check whether data-structures and proofs are suboptimal, spurring code refactoring. Concretely, when developing Frex, we noticed proofs with loops: multi-step derivations that start and end in the same term. Such loops come from internal data structures that optimise simplifier-development effort, but insert semantically irrelevant subterms that can be simplified away. Frex implements a generic

```
units : {a, b : Nat} -> (0 + (a + 0)) + b + 0 = a + b    agdaEx : ∀ {x y} → (2 + x) + (y + 3) ≡ x + (y + 5)
units = %runElab frexMagic MonoidFrexlet Additive          agdaEx = fragment CSemigroupFrex +-csemigroup
```

Fig. 13. Goal extraction in (a) Idris2 Frex's elaborator reflection script; (b) the **frex Ag**da aug**ment**ation lib.

proof simplifier that automatically removes all such loops. This mechanism suggests further work, developing modules for simplifying these proofs further.

6 Goal Extraction via Reflection

Thus far, our examples have illustrated interaction with Frex using `solve`. The `solve` function provides a similar interface to the simplifiers in (e.g.) Agda's standard library: it takes the fral or frex simplifier, the number of free variables and the abstract syntax of a goal. However, these simplifiers additionally provide ergonomic goal extraction with Agda's *proof reflection* mechanism.

Proof reflection is a metaprogramming paradigm, available in proof assistants and dependently typed programming languages, that supports bi-directional communication between a language and its implementation. The language provides a representation of its terms, operators that construct, manipulate and destruct term representations, and primitives *quote* and *unquote* that respectively *reify* terms into the representation and *reflect* back encoded terms as ordinary terms.

Given mechanisms for querying unsolved proof obligations, proof reflection enables the implementation of verified decision procedures for automatically discharging such obligations without boilerplate [Boutin 1997; Christiansen and Brady 2016]. Coupled with the meta-theoretic properties that dependently typed implementations of decision procedures can enforce (e.g. relative soundness and completeness), reflection-driven interfaces yield easy-to-use tactics with strong guarantees. It is then natural to ask: is it possible to construct an interface to Frex that uses proof reflection to avoid the need to explicitly supply the equation to discharge?

As the example code in Fig. 13 illustrates, using proof reflection to provide such an interface to Frex is possible in both Idris2 and Agda. Rather than designing custom reflection-based drivers for individual simplifiers, we combine proof reflection with Frex's design philosophy of extensibility and common core reuse and provide a single generic metaprogram parameterized by a signature and a model of a presentation. The metaprogram can be instantiated for any algebraic simplifier, built-in or user-defined. Fig. 13 (a) shows the invocation of the Idris2 elaboration script `frexMagic`, and Fig. 13 (b) shows the Agda proof reflection macro `fragment`. Both implementations aim to infer the abstract syntax of the goal equation based on the expected type.

The drivers have no information about the structure of the algebraic signature argument ahead of time. Frex's inductive `Term` representation means that relevant abstract operator names can be extracted from the presentation. However, the process of matching goal fragments against the abstract syntax of the algebraic interpretation is tightly coupled to the language's reflection primitives. Implementing Frex in both Idris2 and Agda allows us to compare differences in behaviour.

The differences between the Idris2 and Agda implementations can be seen by considering the normalisation of arithmetic expressions such as `(x + 1) + y = x + (1 + y)`. In Idris2, the reflected syntax passed to the driver represents the normalized expression `(x + 1) + y = x + S y`. As far as the theory of monoids is concerned, `S y` is an atomic expression and so it is treated as another free variable, distinct from `y`. The Idris2 driver then incorrectly infers the invalid equation `Dyn 0 .+. Sta 1 .+. Dyn 1 =-= Dyn 0 .+. Dyn 2`, and fails to discharge the goal. In contrast, Agda does not normalise quoted expressions before reflection, and so the Agda driver successfully finds the equation, allowing Frex to solve the example. Agda's approach is not always superior: there are

similar examples for which the Agda driver fails. The extent to which the implementation can avoid such pathologies ultimately depends on the engineering effort available to develop heuristics.

As these problems indicate, this rather naive approach to automation requires significant developer resources to deal with edge cases or construct bespoke solutions under simplifying assumptions. In large, mature ecosystems it may be possible to maintain practical heuristics despite these challenges. However, there might be better mechanisms for specifying algebraic contexts, allowing the solver to extract the required information automatically; we suggest such directions in Section 10.

7 Supplementary Evaluation: Usability and Interactive Development

The implementation of FREX is still in its early stages, and offers many opportunities for further engineering work to extend its functionality, expressiveness, ergonomics, and efficiency. However, we have already carried out some small experiments to assess user experience and frexlet developer experience to establish that the approach is feasible, and to identify further directions.

7.1 Quantitative evaluation

Idris2 encourages interactive, type-driven development, so it is important that the checker is responsive when the user modifies the program. Following Nielsen [1993], our Idris2 implementation aims for response times under one second, and we treat a response time of over 10 seconds when type checking a modification to FREX client code as a bug.

For typical small equalities that arise incidentally in dependently typed programs, Frex's performance falls very comfortably within Nielsen's limits. For example, the checking time[2] is under 0.1s for terms of size six or below with the commutative solver and terms of size 14 or below with the non-commutative solver, creating an impression of instantaneous response.

As the term size increases, Frex eventually crosses the one second interactivity threshold. Fig. 14 shows how type-checking times grow with term size and with the number of free variables in a randomly generated term for the commutative and non-commutative monoid solvers. As the figure shows, Frex's type-checking time generally remains below the interactivity threshold up to terms of around size 30, and only exceeds the 10 second threshold (beyond which users' attention is lost) for a few terms of size 45 or above. Our experience with Frex development suggests that the anomalously high checking times for these terms is likely to arise from a performance bottleneck in Idris2's evaluator (Section 8) and that the ongoing development of Idris2 may eventually eliminate the problem, bringing the type-checking time for most terms up to size 60 down to a few seconds.

7.2 Qualitative evaluation

To evaluate FREX's usage, we have reproduced Brady et al.'s [2007] dependently typed representation of binary arithmetic. Brady et al. index binary representations by the natural numbers that they represent, and so the programmer needs to prove arithmetic operations correct. Such proofs interleave insightful equational reasoning steps with rote calculational steps such as the following:

```
c_s + 2*(val_s + ((2 `power` width)*c0)) = ((c_s + val_s) + val_s) + (2*((2 `power` width)*c0))
```

which may be discharged with `solve`. We do not use our reflection capabilities since these kinds of examples, in which the binding-time analysis is challenging, are beyond their reach at the moment. With early FREX implementations the task was arduous due to performance bottlenecks in Idris2 that are now eliminated. The only other significant obstacle we encountered was the usual pain point involved in invoking an algebraic simplifier without a goal extraction mechanism: the need to repeat the equation and its relevant rewriting-context when calling FREX.

---

[2]We use a dated AMD FX-8320 machine with 16GB memory, running Idris 2 version 0.5.1-1011cc616 on Debian Linux.

Fig. 14. Frex monoid simplifiers type-checking times

8 System Design Lessons

Frex uses generic and dependently typed programming techniques extensively, requiring significant type level computation that taxes the capabilities of the host language implementation. In developing Frex in Agda and Idris2 we have eliminated some performance bottlenecks in Idris2's type checker, and learned valuable lessons about practical dependently typed language implementation. We share these lessons here, in the hope that developers of other systems will find them useful.

8.1 Idris2

At its heart, the type checker implements: dynamic pattern unification [Gundry 2013; Miller 1992; Reed 2009], which instantiates implicit arguments; and conversion checking whether two terms evaluate to the same reduct. Both components require an evaluator. Idris2 uses normalisation by evaluation [Berger and Schwichtenberg 1991] with a *syntactic representation* (terms) and a *semantic representation* (values in weak head normal form). The static evaluator is call-by-name and produces a weak head normal form from a term, and Idris2 implements a quotation mechanism which reconstructs a term from a semantic representation of a weak head normal form.

Profiling the Idris2 executable reveals that most performance bottlenecks encountered in developing Frex are caused by the evaluator. We have experimented with alternative implementations of the evaluator that compile terms using Scheme's backend and a glued representation of values [Chapman et al. 2005; Coquand and Dybjer 1997] rather than interpreting terms directly. These implementations give modest performance gains, but in the end the most effective way is to avoid evaluation in the first place! We have therefore also experimented with the following ways to

avoid evaluating terms: preserving subterm sharing, choosing appropriate data representation in unification, and taking advantage of the typical structure of unification and conversion problems.

Preserving Sharing Instantiating implicit arguments in dependently typed programs often leads to significant *sharing* of subterms. For example, `[True, False]` : `Vect 2 Bool` elaborates to the following term: `(::) (S Z) Bool True ((::) Z Bool False (Nil Bool))`, sharing the subexpressions `Z` and `Bool`. As the vector gets longer, sharing increases. Following Kovács [2019], we preserve sharing by introducing a metavariable for *every* implicit argument, inlining only when it is guaranteed that the definition cannot break sharing. Consequently, we inline a metavariable whose definition is itself a metavariable applied to local variables. Otherwise, we do not substitute metavariable solutions into terms at all until they are required for unification or display purposes.

Unification Unification operates on *values*, not *terms*, but sometimes Idris2 needs to postpone a unification problem if it is blocked due to an unsolved metavariable. When the metavariable is solved, Idris2 needs to re-evaluate the terms being unified. Previously, Idris2 stored postponed problems as a pair of (syntactic) terms in an environment, re-evaluated once the blocking metavariable is solved. However, Frex produces some large postponed problems, for which quotation to syntax is expensive. Now, in addition to the evaluator and quotation, we have introduced a *continue* operation, which re-evaluates the metavariable at the head of a blocked value, and avoids unnecessary quotation.

Conversion Checking Types in Frex can be large, and sometimes a unification problem that arises while type checking Frex is postponed due to an unsolved metavariable which blocks evaluation. In this case, we might have a unification problem of the form `f x1 ... xn =?= f y1 ... yn` where the `xi`, `yi` etc may be very large subterms, and the terms unify if they are convertible. If most corresponding terms are equal after evaluation, but one differs, it may take a long time to find the differing subterm which blocks unification, especially since checking the convertibility of subterms involves evaluation. Fortunately, terms in blocked unification problems tend to differ at the heads rather than at deeply nested subterms. Therefore, we always check the heads of the values of corresponding `xi` and `yi` first, postponing the unification problem if any are unequal. This heuristic significantly improves performance, preventing a lot of unnecessary evaluation.

Influence on Language Design and Ecosystem The development of Frex has led to the implementation of a number of desirable language features in Idris2. Many of these have been minor changes to the treatment of implicit arguments and `parameters` blocks. More significantly, Frex makes extensive use of `auto` implicit arguments, which are solved by a search procedure which uses constructors and functions marked as search hints. To assist the development of Frex and improve the readability of its code we have added the ability to mark *local* functions as search hints, allowing us to restrict the scope of the hints and so avoid excessive growth of the search space. Frex is now part of the Idris2 test suite, ensuring that it will remain consistent with any updates to Idris2.

8.2 Agda

Agda is a well-established dependently typed interactive proof environment. Idris2 and Agda and their communities have different goals, leading to subtle implementation differences.

The key differences between the two languages arise from Agda's focus on proving versus Idris2's focus on programming. Idris2 currently uses a single *universe* [Palmgren 1998], allowing `Type` : `Type`, and is hence inconsistent by Girard's paradox. In contrast, Agda's well-developed predicative theory of universes avoids Girard's paradox. Agda also protects users from other logical paradoxes of its more experimental features with its '`--safe`' compiler flag. In the spirit of Hu and Carette [2021] and The Agda Community [2024], we adopt a conservative set of compiler options (`--without-K --safe`). All our definitions are universe-polymorphic. This conservativity

broadens the applicability of Frex in the Agda ecosystem by guaranteeing compatibility with all of Agda's various configurations, and further assures us about the correctness of Frex itself. Corbyn [2021] discusses these ideas in greater detail.

9 Related Work

Within the Rocq ecosystem, an abundance of tactics enable algebraic simplification. Boutin's [1997] ring and field tactics[3] let programmers discharge proof obligations involving (and requiring!) addition, multiplication, and division operations. Strub's [2010] CoqMT extends Rocq's Calculus of Inductive Constructions, allowing users to extend the conversion rule with arbitrary decision procedures for first order theories (e.g. Presburger arithmetic). To ensure preservation of good meta-theoretical properties, Strub extends only term level conversion. Implementations of Hilbert's Nullstellensatz theorem (Harrison's [2007] in HOL Light and Pottier's [2008] in Rocq) help users discharge proof obligations involving polynomial equalities on a commutative integral domain.

Rocq's SETOID_REWRITE is an advanced tactic library for setoid rewriting.[4] Disregarding the difference between the direct manipulation of proof-terms in Idris2 and the tactic-based manipulation in Rocq, SETOID_REWRITE provides abstractions for manipulating parameterised relations (covariant and contravariant), and users can register setoids of interest and custom 'morphisms' — horn-like equational clauses — with the library. The various tactics in the library apply these user-defined axioms to the goal. Users may also register tactics, and the library includes an expressive collection of term-traversal primitives (climbing up and down the syntax tree, repeating sub-tactics, and so on). While SETOID_REWRITE does not deal with algebraic simplification directly, it may help in generalising equality-based simplifiers to setoid-based simplifiers. In comparison, Frex's setoid reasoning is minimal, implementing only the necessary features for the library.

In Idris1, Slama and Brady [2017] and Slama [2018] implement a hierarchy of rewriting procedures for algebraic structures of increasing complexity. Though the procedures' completeness is not enforced by type as in Frex, these simplifiers are based on a Knuth-Bendix resolution of critical pairs, and so are likely to be complete. Frex also investigates a hierarchy of rewriting procedures, but: (1) frexlets are complete by construction, (2) Frex is based on normalisation-by-evaluation (like Boutin's tactic, and unlike Slama and Brady's), and (3) Frex is extensible, with support for sufficiently motivated users to extend the library with bespoke solvers.

Normalisation-by-evaluation is an established technique for simplifying terms in a concrete equational theory, often involving function types. One compelling example is Allais et al.'s [2013] work, which demonstrates by a careful model construction that the equational theory decided by normalisation-by-evaluation can be enriched with additional rules. They implement a simply typed language internalising the functorial and fusion laws for list fold, map, and append and prove their construction sound and complete with respect to the extended equational theory.

In Agda, Cockx's [2020] and Cockx et al.'s [2021] '--rewriting' flag allows users to enrich the existing reduction relation with new rules. Their implementation goes beyond Allais et al.'s: it may restart stuck computations. They leave to future work the soundness of user-provided reduction rules, i.e. ensuring rules neither introduce nontermination nor break canonicity. Unlike our approach, neither Allais et al's nor Cockx et al's technique currently supports commutativity.

Existing formalisations for specific algebraic structures (monoids, groups, rings, actions, etc.) abound [Hu and Carette 2021; Mahboubi and Tassi 2022; The Agda Community 2024; The Lean Community 2025]. Frex is set apart from these by formalising the fragments of universal algebra needed for its architecture. However, in Frex we do not aspire to formalise universal algebra for its

---

[3]See the Rocq documentation: https://rocq-prover.org/distrib/current/refman/addendum/ring.html .

[4]See the Rocq documentation: https://rocq-prover.org/refman/addendum/generalized-rewriting.html .

own sake. Formalising more complete fragments of the theory is an active area of research, with recent contributions by Gunther et al. [2018] and Abel [2021] in Agda. Carette et al. [2020] generate Agda code for a comprehensive collection of multi-sorted algebraic theories and their associated machinery via a pre-processing phase from a much smaller description. Fiore and Szamozvancev [2022] similarly generate definitions for the significantly more general second-order abstract syntax in Agda via a preprocessing step, while formalising more of the meta-theory as library code. We consider it an open problem in this domain to include the concise information as first-class data in the meta-language while nonetheless enjoy the full ergonomics of hand-written, inlined, definitions.

Agda's category theory library [Hu and Carette 2021] uses setoids extensively. Each category includes both its homsets and their dualisation. This choice allows for the operation 'take the opposite category' to be judgementally involutive. In turn, one can define dualisation of a property purely formally, e.g. coproducts can be defined as products in the opposite category. Unfortunately, this design makes essential use of η-equality on records, yet unsupported by Idris2.

The Meta-F⋆ language [Martínez et al. 2019] provides normalisation tactics for commutative monoids and semi-rings through its metaprogramming facilities. Frex's usage resembles these tactics' usage, and we hope a Frex port to F⋆ will make use of F⋆'s metaprogramming facilities to reduce some syntactic noise during goal extraction.

10 Conclusions and Further Work

We have presented a novel, mathematically structured, design for algebraic simplification suites that guarantees sound and complete simplification, even of user-defined simplifiers. Preliminary evaluation shows that, despite a high level of abstraction, the resulting library is responsive, with comparable functionality to other libraries, in a combination of features that no single existing library provides. Frex's unique design—the frex and the fral—offer new prospects and questions.

Several computational type theories support extensional function types and quotients, such as observational type theory [Altenkirch et al. 2007; Pujet et al. 2025; Pujet and Tabareau 2022, 2023] and cubical type theory [Cohen et al. 2018; Vezzosi et al. 2019]. Extensional function types and quotient types are also typical reasons for using setoids. So it is natural to consider the tradeoffs for implementing a similar library in a type theory with extensional function types and quotient types.

Even in such rich type theories, we would still want to implement setoid algebras. Setoid algebras let us implement intensional aspects such as proof printing, certification, and simplification, which we view as useful functionality of a simplification library. Quotient types cannot replace this functionality: eliminators of the quotient must respect the quotient.

However, a richer type theory may make other aspects of the library easier. For example, the commutative monoid fral and frex use the power of an algebra by a setoid. If we worked in a type theory with function extensionality, we could use the power of an algebra by a set and do away with some of the setoid machinery involved. A more speculative direction involves using the improvements in the type theory to implement category theoretic abstractions (presheaf categories, ends and coends, 2-functors) which may allow us to 'teach' Idris not just universal algebra, but categorical algebra. We expand on this idea further below.

Cubical type theories offer another potential implementation strategy through non-trivial identity types. Our Agda implementation (Fragment) is compatible with Cubical Agda, as it type-checks with the pragmas `--safe` and `--without-K`. Cubical identity types turn types into untruncated setoids, so our machinery applies unchanged. However, cubical identity types cannot replace setoids. Extracting intensional information is desirable and relevant for cubical type theory too, so we will still implement setoid algebras. So cubical type theory will not improve Frex's core design.

Implementing some of the internal data-structures may be simplified by a more powerful type theory. For example, we can implement the power of a model by an n-element set either as a

function or as an n-vector. It might be possible to use univalence to ease transporting structure between these two different implementations. However, using cubical abstractions—rather than merely being compatible with a cubical type-theory—is a lot more sophisticated. Such transport might lead to unexpected challenges, and would require further experimentation.

Yallop et al.'s [2018] partial evaluators include additional frexlets (e.g. abelian groups, semirings, distributive lattices). We plan to follow suit and port the remaining simplifiers, then conduct larger evaluation and comparison studies. The main challenge is that, unlike Yallop et al., we must mechanising proofs showing these frexlets complete, which is costly. One elegant motivation for including more simplifiers is the following. The frex generalises the 'ring of polynomials over a ring' to that of an algebra of polynomials over an algebra. By porting Yallop et al.'s family of representations, we will fully realise this generalisation.

Our experiment with reflection-based goal extraction as well as the reflection-based interfaces of existing solvers show that with enough engineering efforts, library designers can extract the goal equation from the goal type. However, software engineers for dependently typed languages are a scarce resource. We plan to explore other principled approaches. In practice, when invoked inside a chain of equational steps, the goal equation already appears in the source-code, albeit in a context. Programmers seem willing to type the goal equation once, since it documents the reasoning steps, but seem unhappy to do so *twice*. Generic programming with holes[5] may help.

Another promising direction is *bootstrapping* of the FREX library using simplifier certification. Bootstrapping might start with a hierarchy of inefficient simplifiers that are easy to implement. Next, these simplifiers may then be used to develop a hierarchy of more efficient simplifiers and proof-simplifiers. Finally, the certification mechanism can extract proofs to complete the bootstrap.

We would also like to extend FREX's design beyond algebraic structures. More general notions of theories abound: multisorted, second-order/parameterised, and essentially algebraic. Supporting these may allow FREX to cover much more complex situations, such as decision procedures for first order theories (e.g. Presburger arithmetic, cf. Strub's [2010] CoqMT) normalisation-by-evaluation for fusion laws [Allais et al. 2013], and equational manipulation of big-operators [Bertot et al. 2008; Lau 2017; Markert 2015]. Note that FREX can already deal with big-operators such as `sum` so long as the argument list is a concrete collection of constants and variables such as `sum [2, x]`. We only need the more sophisticated theories when the length of the lists is abstract.

FREX uses many category-theoretic concepts, but the library itself is oblivious to category theory. We hope that making use of a rich category theory library like Hu and Carette's [2021] `agda-categories` might lead to a sleeker and even more modular FREX implementation. More specifically, we plan to explore a general treatment of involutive algebras following Jacobs [2021], and Power's *distributive tensor* of equational theories [Hyland and Power 2006; Power 2005] for a uniform treatment of semi-ring varieties. Instantiating each of the 6 semi-group varieties makes it possible to cover each instance of the following combinations:

$$\begin{pmatrix} \{\text{ordinary}\} \times \{\text{ordinary, involutive, non-reversing involutive}\} \\ \cup\, \{\text{commutative}\} \times \{\text{ordinary, involutive}\} \end{pmatrix} \times \begin{Bmatrix} \text{commutative} \\ \text{semigroup, monoid, group} \end{Bmatrix}$$

and modularly construct $(2 + 3) \times 3 = 15$ semi-ring varieties, including rings and semirings. As this example shows, this kind of modular treatment can provide a multiplicative development boost.

---

[5]See Brad Hardy's Agda-Holes library: https://github.com/bch29/agda-holes .

A   Module Structure of Frex

We summarises the implementation's code modules and their relationship to this article:

- *Frex* (§3–§4.4): core definitions
  - *Signature* (§3): operations & arities
  - *Algebra* (§3,§4.1): algebraic structures and terms, homomorphisms
  - *Presentation* (§3): axioms, equational theories
  - *Axiom* (§3): common axiom schemes
  - *Model* (§3): axiom-validating algebras
  - *Powers* (§4.3): parameterised algebras
  - *Free* (§3): simplification in all algebras
    - *Definition* (§4.1): universal property
    - *Construction* (§5): a non-effective quotient construction used for extraction, printing, and certification
      - *ByFrex* (§4.5): reuse a frex simplifier to define a fral simplifier
      - *Linear* (§5.2–§5.3): generic proof simplification and printing
      - *Idris* (§5): generic certification
  - *Coproduct* (§4.4): universal property
  - *Frex* (§4.1–§4.4): universal property, reuse coproduct and fral simplifier to define a frex simplifier
    - *Construction* (§5): non-effective quotient construction used for extraction, printing, and certification
  - *Lemma* (§5): auxiliary representation for auxiliary lemmata discharged by fral simplifiers, printed, or certified
  - *Magic* (§6): generic reflection code for ergonomic invocation

- *Frexlet.Monoid*: modules concerning varieties of monoids and their simplifiers
  - *Theory* (§3): signature, axioms, pretty printing for the theory of ordinary monoids
  - *Notation* (§3): shared infix notation (additive and multiplicative) for monoid varieties
  - *Frex* (Fig. 1): frex simplifier for monoids
  - *Free* (§4.5): fral simplifier, reuses frex simplifier
  - *Nat* (§3): additive and multiplicative monoid structure of the natural numbers
  - *Pair*: types with the cartesian product as a proof-relevant monoid structure
  - *List*: monoid structure of lists with catenation
  - *Commutative*: commutative monoids modules
    - *Theory*: commutativity axiom
    - *Free* (§4.1:) fral simplifier
    - *NatSemiLinear* (§4.1): auxiliary definitions for fral simplifier
    - *Frex* (§4.4): simplifier, reuses fral via coproducts
    - *Coproduct* (§4.4): coproduct of commutative monoids
    - *Nat*: addition and multiplication of naturals
  - *Involutive*: modules concerning monoids equipped with an involution
    - *Theory* (§3): signature and axioms
    - *Free*: simplifier, reuses frex simplifier
    - *Frex* (§4.6): simplifier, reuses monoid frex
    - *List* (§3): involutive monoid structure of list reversal

B   Extensional Function and Quotient Setoids

Figure 15a defines the quotient of a type by a function `q`, taking two elements to be equal when their images under the function `q` are equal, and the setoid of homomorphisms between two setoids together with extensional equality. This example also demonstrates Idris2's local definitions (lines 4–6, e.g.), possibly with quantities, named-argument function calls (lines 8–15, e.g.), application operator `$`, and anonymous functions (lines 9–10, e.g.). Idris2, like Haskell, implicitly quantifies (with quantity `0`) over unbound variables in type-declarations such as the type `a` in `Quotient`. These underscores mean elaboration must fill-in the blanks uniquely using unification.

Figure 15b presents a setoid over `n`-length vectors over a given setoid. The vector functorial action `VectMap` has a setoid homomorphism structure between the two setoids of homomorphisms: (1) `map f.H` is a homomorphism (lines 19–25), and that (2) it maps extensionally equal homomorphisms to extensionally equal homomorphisms (26–30). These proofs use Idris2's equational reasoning

```
Quotient : (b : Setoid) -> (a -> U b)
  -> Setoid
Quotient b q = MkSetoid a $
  let 0 relation : a -> a -> Type
      relation x y =
        b.equivalence.relation (q x) (q y)
  in MkEquivalence
    { relation = relation
    , reflexive = \x      =>
        b.equivalence.reflexive  (q x)
    , symmetric = \x,y    =>
        b.equivalence.symmetric  (q x) (q y)
    , transitive = \x,y,z =>
        b.equivalence.transitive
          (q x) (q y) (q z)
    }

(~~>) : (a,b : Setoid) -> Setoid
(~~>) a b = MkSetoid (a ~> b) $
  let 0 relation : (f, g : a ~> b) -> Type
      relation f g = (x : U a) ->
        b.equivalence.relation (f.H x) (g.H x)
  in MkEquivalence
  { relation
  , reflexive = \f,v       =>
      b.equivalence.reflexive (f.H v)
  , symmetric = \f,g,prf,w =>
      b.equivalence.symmetric _ _ (prf w)
  , transitive = \f,g,h,f_eq_g, g_eq_h, q =>
      b.equivalence.transitive _ _ _
        (f_eq_g q) (g_eq_h q)
  }
```

```
1 0 (.VectEquality):(a : Setoid) -> Rel (Vect n (U a))
2 a.VectEquality xs ys = (i : Fin n) ->
3   a.equivalence.relation (index i xs) (index i ys)
4 VectSetoid : (n : Nat) -> (a : Setoid) -> Setoid
5 VectSetoid n a = MkSetoid (Vect n (U a))
6   $ MkEquivalence
7   { relation    = (.VectEquality) a
8   , reflexive = \xs                , i =>
9       a.equivalence.reflexive _
10  , symmetric = \xs,ys, prf         , i =>
11      a.equivalence.symmetric _ _   (prf  i)
12  , transitive = \xs, ys, zs, prf1, prf2, i =>
13      a.equivalence.transitive _ _ _ (prf1 i)(prf2 i)
14  }
15 VectMap : {a, b : Setoid} -> (a ~~> b) ~>
16   (VectSetoid n a ~~> VectSetoid n b)
17 VectMap = MkSetoidHomomorphism
18  (\f => MkSetoidHomomorphism
19    (\xs => map f.H xs)
20  $ \xs, ys, prf, i  => CalcWith b $
21  |~ index i (map f.H xs)
22  ~~ f.H(index i xs)
23                      .=.(indexNaturality _ _ _)
24  ~~ f.H(index i ys)...(f.homomorphic _ _ $ prf i)
25  ~~ index i (map f.H ys)
26                      .=<(indexNaturality _ _ _))
27 $ \f,g,prf,xs,i => CalcWith b $
28 |~ index i (map f.H xs)
29 ~~ f.H (index i xs) .=.(indexNaturality _ _ _)
30 ~~ g.H (index i xs) ...(prf _)
31 ~~ index i (map g.H xs) .=<(indexNaturality _ _ _)
```

Fig. 15. (a) Quotient, function-space, and (b) vector setoids (top) and a higher-order homomorphism (bottom)

notation for setoids (lines 20–25 and 27–30), a deeply embedded chain of equational steps. Each step (~~) appeals to transitivity, and requires a justification. The last two dots in the thought bubble operator (...) modify the reason: plain usage (line 23) appeals to a setoid equivalence; an equals in the middle dot, e.g. (.=.), appeals to reflexivity via propositional equality (lines 22, 25, 28, 30); and a comparison symbol in the end, e.g. (.=<), appeals to symmetry (lines 25, 30).

C Proof Printing and Certification

Figure 16 presents the layered representation of linear derivations.

Fig. 17 shows an automatically extracted proof for the equation $(x \bullet 3) \bullet 2 = 5 \bullet x$ in the additive monoid structure (Nat, 0, (+)). The extracted proof has 24 steps — far from the shortest proof possible. Extraction removes reflexivity and transitivity steps, and the pointed bracket tells whether the step uses the axiom directly (angle points right) or using symmetry (angle points left). Square brackets mean appealing to congruence, where the context is the congruence's context, and the term in the

```
data RTList : Rel a -> Rel a where
  Nil  : RTList r x x
  (::) : {0 r : Rel a} -> {y : a}
         -> r x y -> RTList r y z
         -> RTList r x z
```

(a) reflexive-transitive closure

```
data Symmetrise : Rel a -> Rel a where
  Fwd : {0 r : Rel a} -> r x y
         -> Symmetrise r x y
  Bwd : {0 r : Rel a} -> r x y
         -> Symmetrise r y x
```

(b) symmetric closure

```
Derivation : (p : Presentation)
  -> (a : PresetoidAlgebra
          p.signature)
  -> Rel (U a)
Derivation p a
  = RTList       -- Reflexive, Transitive
  $ Symmetrise -- Symmetric
                 -- Congruence
  $ Locate p.signature a.algebra
  $ Step p a   -- Axiomatic steps
```

(e) linear derivations

```
data Locate : (sig : Signature) -> (a : Algebra sig) ->
              Rel (U a) -> Rel (U a) where
  ||| We prove the equality by invoking a rule at the
  ||| toplevel
  Here : {0 r : Rel (U a)} -> r x y
              -> Locate sig a r x y
  ||| We focus on a subterm `lhs` that may appear in
  ||| multiple locations and rewrite it to `rhs` using a
  ||| specific rule.
  Cong : {0 r : Rel (U a)} ->
         (t : Term sig (Maybe (U a))) ->
         {lhs, rhs : U a} -> r lhs rhs ->
         Locate sig a r (plug a t lhs) (plug a t rhs)
```

(c) unary congruence closure

```
data Step : (pres : Presentation)
            -> (a : PresetoidAlgebra pres.signature)
            -> Rel (U a) where
  Include : {x, y : U a} -> a.relation x y
            -> Step pres a x y
  ByAxiom : {0 a : PresetoidAlgebra pres.signature}
            -> (eq : Axiom pres)
            -> (env : Fin (pres.axiom eq).support -> U a)
            -> Step pres a
               (a .bindTerm (pres.axiom eq).lhs env)
               (a .bindTerm (pres.axiom eq).rhs env)
```

(d) axiomatic steps

Fig. 16. Layered (a–d) representation of linear derivations (e)

hole is the equation's LHS. Fig. 18 shows an automatically extracted certificate for the equation $0 + (x + 0) + 0 = x$ in a generic monoid $m = (U m, 01, (.+.))$. The certificate is generated inside a module that parameterises over the generic monoid m and introduces the various notations and reasoning functions.

$$(x \bullet 3) \bullet [2] \underset{\langle \text{ Right neutrality }]}{\overset{\langle \text{ Left neutrality }]}{=}} (x \bullet 3) \bullet 2 \bullet [\varepsilon] \overset{\downarrow}{=} (x \bullet [3]) \bullet 2 \bullet \varepsilon \bullet \varepsilon \underset{\langle \text{ Right neutrality }]}{=} (x \bullet 3 \bullet [\varepsilon]) \bullet 2 \bullet \varepsilon \bullet \varepsilon \overset{\langle \text{ Left neutrality }]}{\underset{\langle \text{ Right neutrality }]}{\overset{\downarrow}{=}}} ([x] \bullet 3 \bullet \varepsilon \bullet \varepsilon) \bullet 2 \bullet \varepsilon \bullet \varepsilon =$$

$$([x \bullet \varepsilon] \bullet 3 \bullet \varepsilon \bullet \varepsilon) \bullet 2 \bullet \varepsilon \bullet \varepsilon \overset{\langle \text{ Right neutrality }]}{=} ([[(x \bullet \varepsilon) \bullet \varepsilon] \bullet 3 \bullet \varepsilon \bullet \varepsilon) \bullet 2 \bullet \varepsilon \bullet \varepsilon \underset{\langle \text{ Left neutrality }]}{=} ((([\varepsilon] \bullet (x \bullet \varepsilon) \bullet \varepsilon) \bullet 3 \bullet \varepsilon \bullet \varepsilon) \bullet 2 \bullet \varepsilon \bullet \varepsilon \overset{[\text{ Evaluate }\rangle}{\overset{\downarrow}{=}}$$

$$[(0 \bullet (x \bullet \varepsilon) \bullet \varepsilon) \bullet 3 \bullet \varepsilon \bullet \varepsilon] \bullet 2 \bullet \varepsilon \bullet \varepsilon \underset{\langle \text{ Associativity }]}{=} (0 \bullet [((x \bullet \varepsilon) \bullet \varepsilon) \bullet 3 \bullet \varepsilon \bullet \varepsilon]) \bullet 2 \bullet \varepsilon \bullet \varepsilon \overset{[\text{ Associativity }\rangle}{\overset{\downarrow}{=}} (0 \bullet [((x \bullet \varepsilon) \bullet \varepsilon) \bullet 3] \bullet \varepsilon \bullet \varepsilon) \bullet 2 \bullet \varepsilon \bullet \varepsilon \overset{[\text{ Commutativity }\rangle}{=}$$

$$(0 \bullet [(3 \bullet (x \bullet \varepsilon) \bullet \varepsilon) \bullet \varepsilon \bullet \varepsilon]) \bullet 2 \bullet \varepsilon \bullet \varepsilon \overset{\langle \text{ Associativity }]}{=} [0 \bullet 3 \bullet ((x \bullet \varepsilon) \bullet \varepsilon) \bullet \varepsilon \bullet \varepsilon] \bullet 2 \bullet \varepsilon \bullet \varepsilon \underset{[\text{ Associativity }\rangle}{=} ((0 \bullet 3) \bullet ((x \bullet \varepsilon) \bullet \varepsilon) \bullet [\varepsilon \bullet \varepsilon]) \bullet 2 \bullet \varepsilon \bullet \varepsilon \overset{[\text{ Left neutrality }\rangle}{\overset{\downarrow}{=}}$$

$$((0 \bullet 3) \bullet [(x \bullet \varepsilon) \bullet \varepsilon] \bullet \varepsilon) \bullet 2 \bullet \varepsilon \bullet \varepsilon \underset{\langle \text{ Associativity }]}{=} ((0 \bullet 3) \bullet (x \bullet [\varepsilon \bullet \varepsilon]) \bullet \varepsilon) \bullet 2 \bullet \varepsilon \bullet \varepsilon \overset{[\text{ Left neutrality }\rangle}{\overset{\downarrow}{=}} ([0 \bullet 3] \bullet (x \bullet \varepsilon) \bullet \varepsilon) \bullet 2 \bullet \varepsilon \bullet \varepsilon \overset{[\text{ Evaluate }\rangle}{=}$$

$$(3 \bullet (x \bullet \varepsilon) \bullet \varepsilon) \bullet 2 \bullet \varepsilon \bullet \varepsilon \underset{[\text{ Associativity }\rangle}{\overset{\langle \text{ Associativity }]}{=}} 3 \bullet [((x \bullet \varepsilon) \bullet \varepsilon) \bullet 2 \bullet \varepsilon \bullet \varepsilon] \overset{\downarrow}{=} 3 \bullet [((x \bullet \varepsilon) \bullet \varepsilon) \bullet 2] \bullet \varepsilon \bullet \varepsilon \overset{[\text{ Commutativity }\rangle}{=}$$

$$3 \bullet [(2 \bullet (x \bullet \varepsilon) \bullet \varepsilon) \bullet \varepsilon \bullet \varepsilon] \underset{\langle \text{ Associativity }]}{=} 3 \bullet 2 \bullet ((x \bullet \varepsilon) \bullet \varepsilon) \bullet \varepsilon \bullet \varepsilon \overset{[\text{ Associativity }\rangle}{\overset{\downarrow}{=}} (3 \bullet 2) \bullet ((x \bullet \varepsilon) \bullet \varepsilon) \bullet [\varepsilon \bullet \varepsilon] \overset{[\text{ Left neutrality }\rangle}{=}$$

$$(3 \bullet 2) \bullet [(x \bullet \varepsilon) \bullet \varepsilon] \bullet \varepsilon \underset{[\text{ Left neutrality }\rangle}{\overset{\langle \text{ Associativity }]}{=}} (3 \bullet 2) \bullet (x \bullet [\varepsilon \bullet \varepsilon]) \bullet \varepsilon \overset{[\text{ Evaluate }\rangle}{=} [3 \bullet 2] \bullet (x \bullet \varepsilon) \bullet \varepsilon \overset{[\text{ Right neutrality }\rangle}{\underset{[\text{ Right neutrality }\rangle}{\overset{\downarrow}{=}}} 5 \bullet [(x \bullet \varepsilon) \bullet \varepsilon] = 5 \bullet [x \bullet \varepsilon] \overset{\downarrow}{=} 5 \bullet x$$

Fig. 17. FREX-extracted proof of $(x \bullet 3) \bullet 2 = 5 \bullet x$ in the additive monoid over Nat

```
units : (x : U m) -> O1 .+. (x .+. O1) .+. O1 =~= x
units x = CalcWith (cast m) $
  |~ O1 .+. (x .+. O1) .+. O1
  ~~ O1 .+. (O1 .+. x .+. O1) .+. O1
    ..<( Cong (\ focus => O2 :+: (focus :+: O2) :+: O2) $ lftNeutrality x )
  ~~ O1 .+. (O1 .+. (x .+. O1)) .+. O1
    ..<( Cong (\ focus => O2 :+: focus :+: O2) $ associativity O1 x O1 )
  ~~ O1 .+. O1 .+. (x .+. O1) .+. O1
    ...( Cong (\ focus => focus :+: O2) $ associativity O1 O1 (x .+. O1) )
  ~~ O1 .+. O1 .+. x .+. O1 .+. O1
    ...( Cong (\ focus => focus :+: O2) $ associativity (O1 .+. O1) x O1 )
  ~~ O1 .+. x .+. O1 .+. O1
    ...( Cong (\ focus => focus :+: Val x :+: O2 :+: O2) $ lftNeutrality O1 )
  ~~ O1 .+. x .+. (O1 .+. O1)
    ..<( associativity (O1 .+. x) O1 O1 )
  ~~ O1 .+. x .+. O1
    ...( Cong (\ focus => O2 :+: Val x :+: focus) $ lftNeutrality O1 )
  ~~ O1 .+. x
    ...( rgtNeutrality (O1 .+. x) )
  ~~ x
    ...( lftNeutrality x )
```

Fig. 18. Frex-certificate for the of $0 + (x + 0) + 0 = x$ in a generic monoid m

Acknowledgments

References

Andreas Abel. 2021. Birkhoff's Completeness Theorem for Multi-Sorted Algebras Formalized in Agda. *CoRR* abs/2111.07936 (2021). arXiv:2111.07936 https://arxiv.org/abs/2111.07936

Andreas Abel, Klaus Aehlig, and Peter Dybjer. 2007a. Normalization by Evaluation for Martin-Löf Type Theory with One Universe. *Electronic Notes in Theoretical Computer Science* 173 (2007), 17–39. https://doi.org/10.1016/j.entcs.2007.02.025 Proceedings of the 23rd Conference on the Mathematical Foundations of Programming Semantics (MFPS XXIII).

Andreas Abel, Thierry Coquand, and Peter Dybjer. 2007b. Normalization by Evaluation for Martin-Löf Type Theory with Typed Equality Judgements. In *22nd Annual IEEE Symposium on Logic in Computer Science (LICS 2007)*. 3–12. https://doi.org/10.1109/LICS.2007.33

Andreas Abel, Andrea Vezzosi, and Théo Winterhalter. 2017. Normalization by evaluation for sized dependent types. *Proc. ACM Program. Lang.* 1, ICFP (2017), 33:1–33:30. https://doi.org/10.1145/3110277

Guillaume Allais, Edwin Brady, Ohad Kammar, and Jeremy Yallop. 2020. Frex: indexing modulo equations with free extensions. (2020). The 5th ACM SIGPLAN International Workshop on Type-Driven Development (TyDe'2020).

Guillaume Allais, Conor McBride, and Pierre Boutillier. 2013. New equations for neutral terms: a sound and complete decision procedure, formalized. In *Proceedings of the 2013 ACM SIGPLAN workshop on Dependently-typed programming, DTP@ICFP 2013, Boston, Massachusetts, USA, September 24, 2013*, Stephanie Weirich (Ed.). ACM, 13–24. https://doi.org/10.1145/2502409.2502411

Thorsten Altenkirch, Paolo Capriotti, Gabe Dijkstra, Nicolai Kraus, and Fredrik Nordvall Forsberg. 2018. Quotient inductive-inductive types. In *International Conference on Foundations of Software Science and Computation Structures*. Springer International Publishing Cham, 293–310.

Thorsten. Altenkirch, Peter. Dybjer, Martin Hofmann, and Philip Scott. 2001. Normalization by evaluation for typed lambda calculus with coproducts. In *Proceedings 16th Annual IEEE Symposium on Logic in Computer Science*. 303–310. https://doi.org/10.1109/LICS.2001.932506

Thorsten Altenkirch, Martin Hofmann, and Thomas Streicher. 1995. Categorical reconstruction of a reduction free normalization proof. In *Category Theory and Computer Science*, David Pitt, David E. Rydeheard, and Peter Johnstone (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 182–199.

Thorsten Altenkirch and Ambrus Kaposi. 2016. Type theory in type theory using quotient inductive types. In *Proceedings of the 43rd Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages* (St. Petersburg, FL, USA) *(POPL '16)*. Association for Computing Machinery, New York, NY, USA, 18–29. https://doi.org/10.1145/2837614.2837638

Thorsten Altenkirch, Conor McBride, and Wouter Swierstra. 2007. Observational equality, now!. In *Proceedings of the 2007 Workshop on Programming Languages Meets Program Verification* (Freiburg, Germany) *(PLPV '07)*. Association for Computing Machinery, New York, NY, USA, 57–68. https://doi.org/10.1145/1292597.1292608

Robert Atkey. 2018. Syntax and Semantics of Quantitative Type Theory. In *Proceedings of the 33rd Annual ACM/IEEE Symposium on Logic in Computer Science, LICS 2018, Oxford, UK, July 09-12, 2018*, Anuj Dawar and Erich Grädel (Eds.). ACM, 56–65. https://doi.org/10.1145/3209108.3209189

Bruno Barras, Benjamin Grégoire, Assia Mahboubi, Laurent Théry, Patrick Loiseleur, and Samuel Boutin. 2021. *The Coq Proof Assistant: Reference Manual. Ring and field: solvers for polynomial and rational equations*. Technical Report. INRIA. Section 3.2.4..

U. Berger and H. Schwichtenberg. 1991. An inverse of the evaluation functional for typed lambda -calculus. In *[1991] Proceedings Sixth Annual IEEE Symposium on Logic in Computer Science*. 203–211. https://doi.org/10.1109/LICS.1991.151645

Yves Bertot, Georges Gonthier, Sidi Ould Biha, and Ioana Pasca. 2008. Canonical Big Operators. In *Theorem Proving in Higher Order Logics*, Otmane Ait Mohamed, César Muñoz, and Sofiène Tahar (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 86–101.

Ilya Beylin and Peter Dybjer. 1996. Extracting a proof of coherence for monoidal categories from a proof of normalization for monoids. In *Types for Proofs and Programs*, Stefano Berardi and Mario Coppo (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 47–61.

Errett Bishop. 1967. *Foundations of Constructive Analysis*. Mcgraw-Hill, New York, NY, USA.

Samuel Boutin. 1997. Using reflection to build efficient and certified decision procedures. *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)* 1281 (1997), 515–529. https://doi.org/10.1007/BFB0014565

Edwin Brady, James McKinna, and Kevin Hammond. 2007. Constructing Correct Circuits: Verification of Functional Aspects of Hardware Specifications with Dependent Types. 159–176. 8th Symposium on Trends in Functional Programming 2007, TFP 2007 ; Conference date: 02-04-2007 Through 04-04-2007.

Stanely Burris and H. P. Sankappanavar. 1981. *A Course in Universal Algebra*. Springer, New York, NY.

Jacques Carette, William M. Farmer, and Yasmine Sharoda. 2020. Leveraging the Information Contained in Theory Presentations. In *Intelligent Computer Mathematics*, Christoph Benzmüller and Bruce Miller (Eds.). Springer International Publishing, Cham, 55–70.

James Chapman, Thorsten Altenkirch, and Conor McBride. 2005. Epigram Reloaded: A Standalone Typechecker for {ETT}. *Trends in Functional Programming* (2005).

David Christiansen and Edwin Brady. 2016. Elaborator Reflection: Extending Idris in Idris. *Proceedings of the 21st ACM SIGPLAN International Conference on Functional Programming* (2016). https://doi.org/10.1145/2951913

Jesper Cockx. 2020. Type Theory Unchained: Extending Agda with User-Defined Rewrite Rules. In *25th International Conference on Types for Proofs and Programs (TYPES 2019) (Leibniz International Proceedings in Informatics (LIPIcs), Vol. 175)*, Marc Bezem and Assia Mahboubi (Eds.). Schloss Dagstuhl–Leibniz-Zentrum für Informatik, Dagstuhl, Germany, 2:1–2:27. https://doi.org/10.4230/LIPIcs.TYPES.2019.2

Jesper Cockx, Nicolas Tabareau, and Théo Winterhalter. 2021. The Taming of the Rew: A Type Theory with Computational Assumptions. *Proc. ACM Program. Lang.* 5, POPL, Article 60 (jan 2021), 29 pages. https://doi.org/10.1145/3434341

Cyril Cohen, Thierry Coquand, Simon Huber, and Anders Mörtberg. 2018. Cubical Type Theory: A Constructive Interpretation of the Univalence Axiom. In *21st International Conference on Types for Proofs and Programs (TYPES 2015) (Leibniz International Proceedings in Informatics (LIPIcs), Vol. 69)*, Tarmo Uustalu (Ed.). Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl, Germany, 5:1–5:34. https://doi.org/10.4230/LIPIcs.TYPES.2015.5

Thierry Coquand and Peter Dybjer. 1997. Intuitionistic Model Constructions and Normalization Proofs. *Math. Struct. Comput. Sci.* 7, 1 (1997), 75–94. https://doi.org/10.1017/S0960129596002150

Nathan Corbyn. 2021. *Proof Synthesis with Free Extensions in Intensional Type Theory*. Technical Report. University of Cambridge. MEng Dissertation.

Djordje Čubrić, Peter Dybjer, and Philip Scott. 1998. Normalization and the Yoneda embedding. *Mathematical Structures in Computer Science* 8, 2 (1998), 153–192. https://doi.org/10.1017/S0960129597002508

Marcelo Fiore and Dmitrij Szamozvancev. 2022. Formal metatheory of second-order abstract syntax. *Proc. ACM Program. Lang.* 6, POPL (2022), 1–29. https://doi.org/10.1145/3498715

Benjamin Grégoire and Assia Mahboubi. 2005. Proving Equalities in a Commutative Ring Done Right in Coq. In *Theorem Proving in Higher Order Logics*, Joe Hurd and Tom Melham (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 98–113.

Jason Gross, Adam Chlipala, and David I. Spivak. 2014. Experience Implementing a Performant Category-Theory Library in Coq. In *Interactive Theorem Proving*, Gerwin Klein and Ruben Gamboa (Eds.). Springer International Publishing, Cham, 275–291.

Adam Gundry. 2013. *Type Inference, Haskell and Dependent Types*. Ph. D. Dissertation. https://personal.cis.strath.ac.uk/adam.gundry/thesis/thesis-2013-07-24.pdf

Emmanuel Gunther, Alejandro Gadea, and Miguel Pagano. 2018. Formalization of Universal Algebra in Agda. *Electronic Notes in Theoretical Computer Science* 338 (2018), 147–166. https://doi.org/10.1016/j.entcs.2018.10.010 The 12th Workshop on Logical and Semantic Frameworks, with Applications (LSFA 2017).

John Harrison. 2007. Automating Elementary Number-Theoretic Proofs Using Gröbner Bases. In *Automated Deduction - CADE-21, 21st International Conference on Automated Deduction, Bremen, Germany, July 17-20, 2007, Proceedings (Lecture Notes in Computer Science, Vol. 4603)*, Frank Pfenning (Ed.). Springer, 51–66. https://doi.org/10.1007/978-3-540-73595-3_5

Martin Hofmann. 1995. *Extensional concepts in intensional type theory*. Ph. D. Dissertation. University of Edinburgh. College of Science and Engineering. School of Informatics. http://hdl.handle.net/1842/399.

Martin Hofmann. 1997. *Extensional constructs in intensional type theory*. Springer.

Jason Z. S. Hu and Jacques Carette. 2021. Formalizing Category Theory in Agda. In *Proceedings of the 10th ACM SIGPLAN International Conference on Certified Programs and Proofs* (Virtual, Denmark) *(CPP 2021)*. Association for Computing

Machinery, New York, NY, USA, 327–342. https://doi.org/10.1145/3437992.3439922

Jason Z. S. Hu, Junyoung Jang, and Brigitte Pientka. 2023. Normalization by evaluation for modal dependent type theory. *J. Funct. Program.* 33 (2023). https://doi.org/10.1017/S0956796823000060

Martin Hyland and John Power. 2006. Discrete Lawvere theories and computational effects. *Theoretical Computer Science* 366, 1 (2006), 144–162. https://doi.org/10.1016/j.tcs.2006.07.007 Algebra and Coalgebra in Computer Science.

Bart Jacobs. 2021. Involutive Categories and Monoids, with a GNS-Correspondence. *Foundations of Physics* 42 (2021), 874–895. Issue 7. https://doi.org/10.1007/s10701-011-9595-7

Ambrus Kaposi, András Kovács, and Thorsten Altenkirch. 2019. Constructing quotient inductive-inductive types. *Proc. ACM Program. Lang.* 3, POPL, Article 2 (Jan. 2019), 24 pages. https://doi.org/10.1145/3290315

Donnacha Oisín Kidney. 2019. *Automatically and Efficiently Illustrating Polynomial Equalities in Agda.* Technical Report. University College Cork. BSc Dissertation.

András Kovács. 2019. Fast Elaboration for Dependent Type Theories. Talk at EU Types WG meeting, 2019.

Stella Lau. 2017. Theory and implementation of a general framework for big operators in Agda. Bachelor's thesis, University of Cambridge.

Saunders Mac Lane. 2010. *Categories for the working mathematician* (2nd. ed., softcover version of original hardcover edition 1998 ed.). Number 5. Springer, New York, NY.

Assia Mahboubi and Enrico Tassi. 2022. *Mathematical Components.* Zenodo. https://doi.org/10.5281/zenodo.7118596

Leonhard Markert. 2015. *Big operators in Agda.* Master's thesis. MSc thesis, University of Cambridge.

Per Martin-Löf. 1975. An Intuitionistic Theory of Types: Predicative Part. In *Logic Colloquium '73*, H.E. Rose and J.C. Shepherdson (Eds.). Studies in Logic and the Foundations of Mathematics, Vol. 80. Elsevier, 73–118. https://doi.org/10.1016/S0049-237X(08)71945-1

Guido Martínez, Danel Ahman, Victor Dumitrescu, Nick Giannarakis, Chris Hawblitzel, Catalin Hritcu, Monal Narasimhamurthy, Zoe Paraskevopoulou, Clément Pit-Claudel, Jonathan Protzenko, Tahina Ramananandro, Aseem Rastogi, and Nikhil Swamy. 2019. Meta-F*: Proof Automation with SMT, Tactics, and Metaprograms. In *28th European Symposium on Programming (ESOP)*. Springer, 30–59. https://doi.org/10.1007/978-3-030-17184-1_2

Conor McBride. 2016. *I Got Plenty o' Nuttin'.* Springer International Publishing, Cham, 207–233. https://doi.org/10.1007/978-3-319-30936-1_12

Dale Miller. 1992. Unification under a mixed prefix. *Journal of Symbolic Computation* (1992). http://www.sciencedirect.com/science/article/pii/074771719290011R

Jakob Nielsen. 1993. *Usability Engineering.* Morgan Kaufmann Publishers Inc., San Francisco, CA, USA.

Erik Palmgren. 1998. On Universes in Type Theory. In *Twenty-Five Years of Constructive Type Theory*, Giovanni Sambin and Jan M. Smith (Eds.). Oxford University Press, Oxford, United Kingdom, Chapter 12, 191–204. https://doi.org/10.1093/oso/9780198501275.003.0012

Loic Pottier. 2008. Connecting Gröbner Bases Programs with Coq to do Proofs in Algebra, Geometry and Arithmetics. In *Proceedings of the LPAR 2008 Workshops, Knowledge Exchange: Automated Provers and Proof Assistants, and the 7th International Workshop on the Implementation of Logics, Doha, Qatar, November 22, 2008 (CEUR Workshop Proceedings, Vol. 418)*, Piotr Rudnicki, Geoff Sutcliffe, Boris Konev, Renate A. Schmidt, and Stephan Schulz (Eds.). CEUR-WS.org. http://ceur-ws.org/Vol-418/paper5.pdf

John Power. 2005. Discrete Lawvere Theories. In *Algebra and Coalgebra in Computer Science*, José Luiz Fiadeiro, Neil Harman, Markus Roggenbach, and Jan Rutten (Eds.). Springer Berlin Heidelberg, Berlin, Heidelberg, 348–363.

Loïc Pujet, Yann Leray, and Nicolas Tabareau. 2025. Observational Equality Meets CIC. *ACM Trans. Program. Lang. Syst.* 47, 2, Article 6 (April 2025), 35 pages. https://doi.org/10.1145/3719342

Loïc Pujet and Nicolas Tabareau. 2022. Observational equality: now for good. *Proc. ACM Program. Lang.* 6, POPL, Article 32 (Jan. 2022), 27 pages. https://doi.org/10.1145/3498693

Loïc Pujet and Nicolas Tabareau. 2023. Impredicative Observational Equality. *Proc. ACM Program. Lang.* 7, POPL, Article 74 (Jan. 2023), 26 pages. https://doi.org/10.1145/3571739

Jason Reed. 2009. Higher-order constraint simplification in dependent type theory. *ACM International Conference Proceeding Series* (2009), 49–56. https://doi.org/10.1145/1577824.1577832

Franck Slama. 2018. *Automatic generation of proof terms in dependently typed programming languages.* Ph. D. Dissertation. http://hdl.handle.net/10023/16451

Franck Slama and Edwin Brady. 2017. Automatically Proving Equivalence by Type-Safe Reflection. In *Intelligent Computer Mathematics*, Herman Geuvers, Matthew England, Osman Hasan, Florian Rabe, and Olaf Teschke (Eds.). Springer International Publishing, Cham, 40–55.

Jonathan Sterling and Carlo Angiuli. 2021. Normalization for Cubical Type Theory. In *2021 36th Annual ACM/IEEE Symposium on Logic in Computer Science (LICS)*. 1–15. https://doi.org/10.1109/LICS52264.2021.9470719

Pierre-Yves Strub. 2010. Coq Modulo Theory. In *Computer Science Logic, 24th International Workshop, CSL 2010, 19th Annual Conference of the EACSL, Brno, Czech Republic, August 23-27, 2010. Proceedings (Lecture Notes in Computer Science,*

Vol. 6247), Anuj Dawar and Helmut Veith (Eds.). Springer, 529–543. https://doi.org/10.1007/978-3-642-15205-4_40

The Agda Community. 2024. *Agda Standard Library*. https://github.com/agda/agda-stdlib

The Lean Community. 2025. *Lean* `mathlib4`. https://github.com/leanprover-community/mathlib4

Atze van der Ploeg and Oleg Kiselyov. 2014. Reflection without remorse: revealing a hidden sequence to speed up monadic reflection. In *Proceedings of the 2014 ACM SIGPLAN symposium on Haskell, Gothenburg, Sweden, September 4-5, 2014*, Wouter Swierstra (Ed.). ACM, 133–144. https://doi.org/10.1145/2633357.2633360

Andrea Vezzosi, Anders Mörtberg, and Andreas Abel. 2019. Cubical agda: a dependently typed programming language with univalence and higher inductive types. *Proc. ACM Program. Lang.* 3, ICFP, Article 87 (July 2019), 29 pages. https://doi.org/10.1145/3341691

Jeremy Yallop, Tamara von Glehn, and Ohad Kammar. 2018. Partially-Static Data as Free Extension of Algebras. *Proc. ACM Program. Lang.* 2, ICFP, Article 100 (July 2018), 30 pages. https://doi.org/10.1145/3236795